

# Day16

## Table of Contents

- [1. Scripts](#)
  - [1.1. compile script](#)
  - [1.2. run script](#)
- [2. MPI Primitive Data Types](#)
- [3. MPI Derived Data Types](#)
- [4. MPI\\_Type\\_contiguous](#)
  - [4.1. Compilation and Execution](#)
- [5. MPI\\_Type\\_vector](#)
  - [5.1. Syntax](#)
  - [5.2. Example Code](#)
  - [5.3. Explanation](#)
  - [5.4. Compilation and Execution](#)
- [6. MPI\\_Type\\_vector Example2](#)
  - [6.1. Compilation and Execution](#)
- [7. MPI\\_Type\\_indexed](#)
  - [7.1. Syntax](#)
  - [7.2. Example Code](#)
  - [7.3. Explanation](#)
  - [7.4. Compilation and Execution](#)
- [8. MPI\\_Type\\_struct](#)
  - [8.1. Syntax \(Deprecated\)](#)
  - [8.2. Syntax \(Current\)](#)
  - [8.3. Example Code](#)
  - [8.4. Explanation](#)
  - [8.5. Compilation and Execution](#)
- [9. MPI\\_Type\\_Struct with different blocklength](#)
  - [9.1. Compilation and Execution](#)
- [10. Reference](#)
  - [10.1. MPI Type contiguous](#)
  - [10.2. MPI Type vector](#)
  - [10.3. MPI Type indexed](#)
  - [10.4. MPI Type struct](#)
- [11. task1](#)
  - [11.1. Compilation and Execution](#)
- [12. task2](#)
  - [12.1. Compilation and Execution](#)

- [13. task2 with MPI\\_Type\\_contiguous](#)
  - [13.1. Compilation and Execution](#)
- [14. task3](#)
  - [14.1. Compilation and Execution](#)
- [15. task4 \(transfer 3rd column to process 1\)](#)
  - [15.1. Compilation and Execution](#)
- [16. Implement vector using indexed](#)
  - [16.1. Compilation and Execution](#)
- [17. Implement contiguous using indexed](#)
  - [17.1. Compilation and Execution](#)

## 1. Scripts

### 1.1. compile script

```
#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cb
#spack load openmpi/c7kvyyq
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

inputFile=$1
outputFile="${1%.*}.out"      # extract the name of the file without extension and adding extension .out
#cmd=`mpicc $inputFile -o $outputFile`
cmd="mpicc $inputFile -o $outputFile"      # running code using MPI
echo "-----"
echo "Command executed: $cmd"
echo "-----"
$cmd

echo "Compilation successful. Check at $outputFile"
echo "-----"
```

### 1.2. run script

```
#!/bin/sh

#source /opt/ohpc/pub/apps/spack/share/spack/setup-env.sh
#spack load gcc/5i5y5cbc
source ~/git/spack/share/spack/setup-env.sh
spack load openmpi

cmd="mpirun -np $2 $1"
```

```

echo "-----"
echo "Command executed: $cmd"
echo "-----"
echo "#####"
echo "#####          OUTPUT          #####"
echo "#####"
echo
mpirun -np $2 $1
echo
echo "#####"
echo "#####          DONE          #####"
echo "#####"

```

## 2. MPI Premitive Data Types

- MPI\_CHAR
- MPI\_WCHAR
- MPI\_SHORT
- MPI\_INT
- MPI\_LONG
- MPI\_LONG\_LONG\_INT
- MPI\_LONG\_LONG
- MPI\_SIGNED\_CHAR
- MPI\_UNSIGNED\_CHAR
- MPI\_UNSIGNED\_SHORT
- MPI\_UNSIGNED\_LONG
- MPI\_UNSIGNED
- MPI\_FLOAT
- MPI\_DOUBLE
- MPI\_LONG\_DOUBLE

## 3. MPI Derived Data Types

- Contiguous
- Vector
- Indexed
- Struct

## 4. MPI\_Type\_contiguous

`MPI\_Type\_contiguous` creates a new MPI datatype that represents a contiguous block of elements. This is useful when you want to send or receive a block of the same datatype as a single message.

```

#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#define N 20

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    const int count = 5;
    int data[N];
    MPI_Datatype contiguous_type;

    // Create a contiguous datatype
    MPI_Type_contiguous(count, MPI_INT, &contiguous_type);
    MPI_Type_commit(&contiguous_type);

    if (rank == 0) {
        // Initialize the data array with some values
        for (int i = 0; i < N; i++) {
            data[i] = i + 1;
        }

        MPI_Send(data, 4, contiguous_type, 1, 0, MPI_COMM_WORLD);
        printf("Process 0 sent data: ");
        for (int i = 0; i < N; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    } else if (rank == 1) {
        MPI_Recv(data, 4, contiguous_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Process 1 received data: ");
        for (int i = 0; i < N; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    }

    MPI_Type_free(&contiguous_type);
    MPI_Finalize();
    return 0;
}

```

## 4.1. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_type_contiguous.c
```

```
-----  
Command executed: mpicc mpi_type_contiguous.c -o mpi_type_contiguous.out  
-----  
Compilation successful. Check at mpi_type_contiguous.out  
-----
```

- Run the program:

```
bash run.sh ./mpi_type_contiguous.out 2
```

```
-----  
Command executed: mpirun -np 2 ./mpi_type_contiguous.out  
-----  
#####  
#####          OUTPUT          #####  
#####  
#####  
  
Process 0 sent data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20  
Process 1 received data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20  
  
#####  
#####          DONE          #####  
#####  
#####
```

In this example, `MPI\_Type\_contiguous` is used to create a contiguous datatype that represents an array of integers. This datatype is then used to send and receive the array between processes.

## 5. MPI\_Type\_vector

`MPI\_Type\_vector` creates a new MPI datatype that represents a pattern of regularly spaced blocks of data. This is useful for sending or receiving non-contiguous data with a regular pattern, such as columns of a matrix or every nth element of an array.

### 5.1. Syntax

```
int MPI_Type_vector(int count, int blocklength, int stride, MPI_Datatype oldtype, MPI_Datatype *newtype);
```

- ``count``: Number of blocks.
- ``blocklength``: Number of elements in each block.
- ``stride``: Number of elements between the start of each block.
- ``oldtype``: Datatype of each element in the block.
- ``newtype``: New datatype representing the vector.

## 5.2. Example Code

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    const int count = 3;          // Number of blocks
    const int blocklength = 3;    // Number of elements in each block
    const int stride = 6;         // Number of elements between the start of each block
    int data[15];                 // Array to send/receive
    MPI_Datatype vector_type;
    MPI_Type_vector(count, blocklength, stride, MPI_INT, &vector_type);
    MPI_Type_commit(&vector_type);
    if (rank == 0) {
        for (int i = 0; i < 15; i++) {
            data[i] = i + 1;
        }

        MPI_Send(data, 1, vector_type, 1, 0, MPI_COMM_WORLD);
        printf("Process 0 sent data: ");
        for (int i = 0; i < 15; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    } else if (rank == 1) {
        for (int i = 0; i < 15; i++) {
            data[i] = 0;
        }

        MPI_Recv(data, 1, vector_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Process 1 received data: ");
        for (int i = 0; i < 15; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    }
}
```

```

    }
    MPI_Type_free(&vector_type);
    MPI_Finalize();
    return 0;
}

```

### 5.3. Explanation

- **Initialization:** Initialize MPI, get the rank and size of the communicator.
- **Datatype Creation:** `MPI\_Type\_vector` creates a new datatype `vector\_type` representing 3 blocks of 1 integer each, with a stride of 5 integers between the start of each block.
- **Process 0:**
  - Initializes the `data` array with values from 1 to 15.
  - Sends the `data` array using the `vector\_type` to process 1.
  - Prints the `data` array.
- **Process 1:**
  - Initializes the `data` array to zero.
  - Receives the data from process 0 into the `data` array using the `vector\_type`.
  - Prints the `data` array after receiving.
- **Datatype Cleanup:** Free the `vector\_type` with `MPI\_Type\_free`.
- **Finalize:** Finalize the MPI environment.

### 5.4. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_type_vector.c
```

```

-----
Command executed: mpicc mpi_type_vector.c -o mpi_type_vector.out
-----
Compilation successful. Check at mpi_type_vector.out
-----

```

- Run the program:

```
bash run.sh ./mpi_type_vector.out 2
```

```

-----
Command executed: mpirun -np 2 ./mpi_type_vector.out
-----

```

```
#####
#####          OUTPUT          #####
#####

Process 0 sent data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15
Process 1 received data: 1 2 3 0 0 0 7 8 9 0 0 0 13 14 15

#####
#####          DONE          #####
#####
```

This example demonstrates how to use `MPI\_Type\_vector` to communicate non-contiguous data with a regular pattern in MPI.

## 6. MPI\_Type\_vector Example2

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    const int count = 3;           // Number of blocks
    const int blocklength = 2;     // Number of elements in each block
    const int stride = 5;          // Number of elements between the start of each block
    MPI_Datatype vector_type;

    // Create a vector datatype
    MPI_Type_vector(count, blocklength, stride, MPI_INT, &vector_type);
    MPI_Type_commit(&vector_type);

    if (rank == 0) {
        int data[15];              // Array to send/receive
        // Initialize the data array with some values
        for (int i = 0; i < 15; i++) {
            data[i] = i + 1;
        }

        MPI_Send(data, 1, vector_type, 1, 0, MPI_COMM_WORLD);
        printf("Process 0 sent data: ");
        for (int i = 0; i < 15; i++) {
```



```

        printf("%d ", data[i]);
    }
    printf("\n");
} else if (rank == 1) {
    int data1[6];

    MPI_Recv(data1, count * blocklength, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received data: ");
    for (int i = 0; i < 6; i++) {
        printf("%d ", data1[i]);
    }
    printf("\n");
}

MPI_Type_free(&vector_type);
MPI_Finalize();
return 0;
}

```

## 6.1. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_type_vector1.c
```

```

-----
Command executed: mpicc mpi_type_vector1.c -o mpi_type_vector1.out
-----
Compilation successful. Check at mpi_type_vector1.out
-----

```

- Run the program:

```
bash run.sh ./mpi_type_vector1.out 2
```

```

-----
Command executed: mpirun -np 2 ./mpi_type_vector1.out
-----
#####
#####              OUTPUT              #####
#####
#####
Process 0 sent data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15
Process 1 received data: 1 2 6 7 11 12
#####

```

```
#####          DONE          #####
#####
```

This example demonstrates how to use ``MPI_Type_vector`` to communicate non-contiguous data with a regular pattern in MPI.

## 7. MPI\_Type\_indexed

``MPI_Type_indexed`` creates a new MPI datatype that represents an irregularly spaced set of blocks of data. This is useful for sending or receiving non-contiguous data with an irregular pattern.

### 7.1. Syntax

```
int MPI_Type_indexed(int count, const int array_of_blocklengths[], const int array_of_displacements[], MPI_Datatype oldtype, MPI_Datatype *newtype)
```

- ``count``: Number of blocks.
- ``array_of_blocklengths``: Array specifying the number of elements in each block.
- ``array_of_displacements``: Array specifying the displacement of each block from the start.
- ``oldtype``: Datatype of each element in the blocks.
- ``newtype``: New datatype representing the indexed pattern.

### 7.2. Example Code

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    const int count = 3;
    int blocklengths[3] = {1, 2, 1};
    int displacements[3] = {0, 3, 7};
    int data[10];
    MPI_Datatype indexed_type;
```

```

// Create an indexed datatype
MPI_Type_indexed(count, blocklengths, displacements, MPI_INT, &indexed_type);
MPI_Type_commit(&indexed_type);

if (rank == 0) {
    // Initialize the data array with some values
    for (int i = 0; i < 10; i++) {
        data[i] = i + 1;
    }

    MPI_Send(data, 1, indexed_type, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent data: ");
    for (int i = 0; i < 10; i++) {
        printf("%d ", data[i]);
    }
    printf("\n");
} else if (rank == 1) {
    // Initialize the data array to zero
    for (int i = 0; i < 10; i++) {
        data[i] = 0;
    }

    MPI_Recv(data, 1, indexed_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received data: ");
    for (int i = 0; i < 10; i++) {
        printf("%d ", data[i]);
    }
    printf("\n");
}

MPI_Type_free(&indexed_type);
MPI_Finalize();
return 0;
}

```

### 7.3. Explanation

- **Initialization:** Initialize MPI, get the rank and size of the communicator.
- **Datatype Creation:**
  - ``blocklengths`` specifies the number of elements in each block: {1, 2, 1}.
  - ``displacements`` specifies the starting indices of each block: {0, 3, 7}.
  - ``MPI_Type_indexed`` creates a new datatype ``indexed_type`` representing these blocks.
- **Process 0:**
  - Initializes the ``data`` array with values from 1 to 10.
  - Sends the ``data`` array using the ``indexed_type`` to process 1.
  - Prints the ``data`` array.
- **Process 1:**
  - Initializes the ``data`` array to zero.

- Receives the data from process 0 into the `data` array using the `indexed\_type`.
- Prints the `data` array after receiving.
- **Datatype Cleanup:** Free the `indexed\_type` with `MPI\_Type\_free`.
- **Finalize:** Finalize the MPI environment.

## 7.4. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_type_indexed.c
```

```
-----
Command executed: mpicc mpi_type_indexed.c -o mpi_type_indexed.out
-----
Compilation successful. Check at mpi_type_indexed.out
-----
```

- Run the program:

```
bash run.sh ./mpi_type_indexed.out 2
```

```
-----
Command executed: mpirun -np 2 ./mpi_type_indexed.out
-----
#####
#####              OUTPUT              #####
#####
#####
Process 0 sent data: 1 2 3 4 5 6 7 8 9 10
Process 1 received data: 1 0 0 4 5 0 0 8 0 0
#####
#####              DONE              #####
#####
```

This example demonstrates how to use `MPI\_Type\_indexed` to communicate non-contiguous data with an irregular pattern in MPI.

## 8. MPI\_Type\_struct

`MPI\_Type\_struct` (now deprecated and replaced by `MPI\_Type\_create\_struct`) allows you to create a new MPI datatype that consists of a sequence of blocks, each with potentially different types and sizes. This is

useful for sending or receiving complex data structures, such as structs in C.

### 8.1. Syntax (Deprecated)

```
int MPI_Type_struct(int count, const int array_of_blocklengths[], const MPI_Aint array_of_displacements[], const MPI_Datatype array_of_types[],
```

### 8.2. Syntax (Current)

```
int MPI_Type_create_struct(int count, const int array_of_blocklengths[], const MPI_Aint array_of_displacements[], const MPI_Datatype array_of_ty
```

- ``count``: Number of blocks.
- ``array_of_blocklengths``: Array specifying the number of elements in each block.
- ``array_of_displacements``: Array specifying the byte displacement of each block from the start.
- ``array_of_types``: Array specifying the datatype of each block.
- ``newtype``: New datatype representing the struct.

### 8.3. Example Code

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

typedef struct {
    int a;
    double b;
    char c;
} my_struct;

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    my_struct data;
    MPI_Datatype struct_type;

    // Create the datatype for my_struct
```

```

int blocklengths[3] = {1, 1, 1};
MPI_Aint displacements[3];
MPI_Datatype types[3] = {MPI_INT, MPI_DOUBLE, MPI_CHAR};

displacements[0] = offsetof(my_struct, a);
displacements[1] = offsetof(my_struct, b);
displacements[2] = offsetof(my_struct, c);

MPI_Type_create_struct(3, blocklengths, displacements, types, &struct_type);
MPI_Type_commit(&struct_type);

if (rank == 0) {
    data.a = 42;
    data.b = 3.14;
    data.c = 'A';

    MPI_Send(&data, 1, struct_type, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent struct: {a = %d, b = %.2f, c = %c}\n", data.a, data.b, data.c);
} else if (rank == 1) {
    MPI_Recv(&data, 1, struct_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received struct: {a = %d, b = %.2f, c = %c}\n", data.a, data.b, data.c);
}

MPI_Type_free(&struct_type);
MPI_Finalize();
return 0;
}

```

## 8.4. Explanation

- **Initialization:** Initialize MPI, get the rank and size of the communicator.
- **Datatype Creation:**
  - `blocklengths` specifies the number of elements in each block: {1, 1, 1}.
  - `displacements` specifies the byte offsets of each block within the struct: calculated using `offsetof`.
  - `types` specifies the datatype of each block: {MPI\_INT, MPI\_DOUBLE, MPI\_CHAR}.
  - `MPI_Type_create_struct` creates a new datatype `struct_type` representing the `my_struct`.
- **Process 0:**
  - Initializes the `data` struct with values.
  - Sends the `data` struct using the `struct_type` to process 1.
  - Prints the `data` struct.
- **Process 1:**
  - Receives the struct from process 0 into the `data` struct using the `struct_type`.
  - Prints the `data` struct after receiving.
- **Datatype Cleanup:** Free the `struct_type` with `MPI_Type_free`.
- **Finalize:** Finalize the MPI environment.

## 8.5. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_type_struct.c
```

```
-----  
Command executed: mpicc mpi_type_struct.c -o mpi_type_struct.out  
-----  
Compilation successful. Check at mpi_type_struct.out  
-----
```

- Run the program:

```
bash run.sh ./mpi_type_struct.out 2
```

```
-----  
Command executed: mpirun -np 2 ./mpi_type_struct.out  
-----  
#####  
#####          OUTPUT          #####  
#####  
#####  
Process 0 sent struct: {a = 42, b = 3.14, c = A}  
Process 1 received struct: {a = 42, b = 3.14, c = A}  
#####  
#####          DONE          #####  
#####
```

This example demonstrates how to use `MPI\_Type\_create\_struct` to communicate complex data structures in MPI.

## 9. MPI Type Struct with different blocklength

```
#include <mpi.h>  
#include <stdio.h>  
#include <stdlib.h>  
#include <stddef.h>  
  
typedef struct {  
    int arr[3];  
    double b;  
    char c;
```

```

} my_struct;

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }

    my_struct data;
    MPI_Datatype struct_type;

    // Create the datatype for my_struct
    int blocklengths[3] = {2, 1, 1}; // Sending part of the array, the double, and the char
    MPI_Aint displacements[3];
    MPI_Datatype types[3] = {MPI_INT, MPI_DOUBLE, MPI_CHAR};

    displacements[0] = offsetof(my_struct, arr);
    displacements[1] = offsetof(my_struct, b);
    displacements[2] = offsetof(my_struct, c);

    MPI_Type_create_struct(3, blocklengths, displacements, types, &struct_type);
    MPI_Type_commit(&struct_type);

    if (rank == 0) {
        data.arr[0] = 1;
        data.arr[1] = 2;
        data.arr[2] = 3;
        data.b = 3.14;
        data.c = 'A';

        MPI_Send(&data, 1, struct_type, 1, 0, MPI_COMM_WORLD);
        printf("Process 0 has struct: {arr = [%d, %d, %d], b = %.2f, c = %c}\n", data.arr[0], data.arr[1], data.arr[2], data.b, data.c);
    } else if (rank == 1) {
        // Initialize the struct to zero
        data.arr[0] = data.arr[1] = data.arr[2] = 0;
        data.b = 0.0;
        data.c = '0';

        MPI_Recv(&data, 1, struct_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Process 1 received struct: {arr = [%d, %d, %d], b = %.2f, c = %c}\n", data.arr[0], data.arr[1], data.arr[2], data.b, data.c);
    }

    MPI_Type_free(&struct_type);
    MPI_Finalize();
    return 0;
}

```



## 9.1. Compilation and Execution

- Compile the program:

```
bash compile.sh mpi_type_struct2.c
```

```
-----  
Command executed: mpicc mpi_type_struct2.c -o mpi_type_struct2.out  
-----  
Compilation successful. Check at mpi_type_struct2.out  
-----
```

- Run the program:

```
bash run.sh ./mpi_type_struct2.out 2
```

```
-----  
Command executed: mpirun -np 2 ./mpi_type_struct2.out  
-----  
#####  
#####          OUTPUT          #####  
#####  
#####  
  
Process 0 has struct: {arr = [1, 2, 3], b = 3.14, c = A}  
Process 1 received struct: {arr = [1, 2, 0], b = 3.14, c = A}  
  
#####  
#####          DONE          #####  
#####
```

This example demonstrates how to use ``MPI_Type_create_struct`` to communicate complex data structures in MPI, specifically how to send parts of an array along with other fields.

## 10. Reference

These images are for your reference. [Link for your reference](#) if you want to learn more about it.

### 10.1. MPI Type contiguous

## MPI\_Type\_contiguous

```
count = 4;  
MPI_Type_contiguous(count, MPI_FLOAT, &rowtype);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

a[4][4]

```
MPI_Send(&a[2][0], 1, rowtype, dest, tag, comm);
```

9.0	10.0	11.0	12.0
-----	------	------	------

1 element of  
rowtype

### 10.2. MPI Type vector

## MPI\_Type\_vector

```
count = 4; blocklength = 1; stride = 4;  
MPI_Type_vector(count, blocklength, stride, MPI_FLOAT,  
                &column_type);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0


a[4][4]

```
MPI_Send(&a[0][1], 1, column_type, dest, tag, comm);
```

2.0	6.0	10.0	14.0
-----	-----	------	------


1 element of  
column\_type

## MPI\_Type\_indexed



```
MPI_Type_indexed(count, blocklengths, displacements, MPI_FLOAT, &indextype);
```

```
MPI_Send(&a, 1, indextype, dest, tag, comm);
```




1 element of  
indextype

# MPI\_Type\_struct

```
typedef struct { float x,y,z,velocity; int n,type; } Particle;  
Particle particles[NELEM];
```

```
MPI_Type_extent(MPI_FLOAT, &extent);
```

```
count = 2; oldtypes[0] = MPI_FLOAT; oldtypes[1] = MPI_INT  
offsets[0] = 0; offsets[1] = 4 * extent;  
blockcounts[0] = 4; blockcounts[1] = 2;
```



```
MPI_Type_struct(count, blockcounts, offsets, oldtypes, &particletype);
```

```
MPI_Send(particles, NELEM, particletype, dest, tag, comm);
```

Sends entire (NELEM) array of particles, each particle being comprised four floats and two integers.

## 11. task1

```
#include <mpi.h>  
#include <stdio.h>  
#include <stdlib.h>  
#define N 5
```

```

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }
    int data[N];

    if (rank == 0) {
        // Initialize the data array with some values
        for (int i = 0; i < N; i++) {
            data[i] = i + 1;
        }

        MPI_Send(data, N, MPI_INT, 1, 0, MPI_COMM_WORLD);
        printf("Process 0 sent data: ");
        for (int i = 0; i < N; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    } else if (rank == 1) {
        MPI_Recv(data, N, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Process 1 received data: ");
        for (int i = 0; i < N; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    }

    MPI_Finalize();
    return 0;
}

```

## 11.1. Compilation and Execution

- Compile the program:

```
bash compile.sh task1.c
```

```
-----  
Command executed: mpicc task1.c -o task1.out  
-----
```

```
Compilation successful. Check at task1.out  
-----
```

- Run the program:

```
bash run.sh ./task1.out 2
```

```
-----  
Command executed: mpirun -np 2 ./task1.out  
-----
```

```
#####  
#####          OUTPUT          #####  
#####
```

```
Process 0 sent data: 1 2 3 4 5  
Process 1 received data: 1 2 3 4 5
```

```
#####  
#####          DONE          #####  
#####
```

## 12. task2

Transfer a matrix of size  $N * N$ .

- Initialize a matrix for both process
- Allow only one process to create whole matrix
- That process will transfer the matrix to another process
- Print the matrix received.

```
#include <mpi.h>  
#include <stdio.h>  
#include <stdlib.h>  
#define N 5  
  
int main(int argc, char** argv) {  
    MPI_Init(&argc, &argv);  
  
    int rank, size;  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
    MPI_Comm_size(MPI_COMM_WORLD, &size);
```

```

if (size < 2) {
    fprintf(stderr, "World size must be greater than 1 for this example\n");
    MPI_Abort(MPI_COMM_WORLD, 1);
}
int data[N][N];

if (rank == 0) {
    int temp = 1;
    // Initialize the data array with some values
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            data[i][j] = temp;
            temp++;
        }
    }

    MPI_Send(data, N * N, MPI_INT, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent data: \n");
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            printf("%d\t", data[i][j]);
        }
        printf("\n");
    }
    printf("\n");
} else if (rank == 1) {
    MPI_Recv(data, N * N, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received data: \n");
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            printf("%d\t", data[i][j]);
        }
        printf("\n");
    }
}

MPI_Finalize();
return 0;
}

```

## 12.1. Compilation and Execution

- Compile the program:

```
bash compile.sh task2.c
```

```
-----
Command executed: mpicc task2.c -o task2.out
-----
```

```
Compilation successful. Check at task2.out
```



- Run the program:

```
bash run.sh ./task2.out 2
```

```
-----
Command executed: mpirun -np 2 ./task2.out
-----
#####
#####              OUTPUT              #####
#####

Process 0 sent data:
 1   2   3   4   5
 6   7   8   9  10
11  12  13  14  15
16  17  18  19  20
21  22  23  24  25

Process 1 received data:
 1   2   3   4   5
 6   7   8   9  10
11  12  13  14  15
16  17  18  19  20
21  22  23  24  25

#####
#####              DONE              #####
#####
```

## 13. task2 with MPI\_Type\_contiguous

Transfer a matrix of size  $N * N$ .

- Initialize a matrix for both process
- Allow only one process to create whole matrix
- That process will transfer the matrix to another process
- Print the matrix received.

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#define N 5

int main(int argc, char** argv) {
```

```

MPI_Init(&argc, &argv);

int rank, size;
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

if (size < 2) {
    fprintf(stderr, "World size must be greater than 1 for this example\n");
    MPI_Abort(MPI_COMM_WORLD, 1);
}
int data[N][N];
MPI_Datatype contiguous_type;
MPI_Type_contiguous(N, MPI_INT, &contiguous_type);
MPI_Type_commit(&contiguous_type);

if (rank == 0) {
    int temp = 1;
    // Initialize the data array with some values
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            data[i][j] = temp;
            temp++;
        }
    }

    MPI_Send(data[3], 2, contiguous_type, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent data: \n");
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            printf("%d\t", data[i][j]);
        }
        printf("\n");
    }
    printf("\n");
} else if (rank == 1) {
    MPI_Recv(data[3], 2, contiguous_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received data: \n");
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            printf("%d\t", data[i][j]);
        }
        printf("\n");
    }
}

MPI_Type_free(&contiguous_type);
MPI_Finalize();
return 0;
}

```

### 13.1. Compilation and Execution

- Compile the program:

```
bash compile.sh task2_contiguous.c
```

```
-----
Command executed: mpicc task2_contiguous.c -o task2_contiguous.out
-----
Compilation successful. Check at task2_contiguous.out
-----
```

- Run the program:

```
bash run.sh ./task2_contiguous.out 2
```

```
-----
Command executed: mpirun -np 2 ./task2_contiguous.out
-----
#####
#####              OUTPUT              #####
#####

Process 0 sent data:
1      2      3      4      5
6      7      8      9     10
11     12     13     14     15
16     17     18     19     20
21     22     23     24     25

Process 1 received data:
49152 0      -1630691416  32767  6
154   0      0          0      0
0     0      0          0      0
16    17     18         19     20
21    22     23         24     25

#####
#####              DONE              #####
#####
```

## 14. task3

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
```

```

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);
    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        fprintf(stderr, "World size must be greater than 1 for this example\n");
        MPI_Abort(MPI_COMM_WORLD, 1);
    }
    const int count = 3;          // Number of blocks
    const int blocklength = 4;    // Number of elements in each block
    const int stride = 6;         // Number of elements between the start of each block
    MPI_Datatype vector_type;
    MPI_Type_vector(count, blocklength, stride, MPI_INT, &vector_type);
    MPI_Type_commit(&vector_type);
    if (rank == 0) {
        int data[20];             // Array to send/receive
        for (int i = 0; i < 20; i++) {
            data[i] = i + 1;
        }
        MPI_Send(data, 1, vector_type, 1, 0, MPI_COMM_WORLD);
        printf("Process 0 sent data: ");
        for (int i = 0; i < 20; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    } else if (rank == 1) {
        int size = count * blocklength;
        int data[size];
        MPI_Recv(data, size, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        printf("Process 1 received data: ");
        for (int i = 0; i < size; i++) {
            printf("%d ", data[i]);
        }
        printf("\n");
    }
    MPI_Type_free(&vector_type);
    MPI_Finalize();
    return 0;
}

```

## 14.1. Compilation and Execution

- Compile the program:

```
bash compile.sh task3.c
```

```
-----
Command executed: mpicc task3.c -o task3.out
```

```
-----  
Compilation successful. Check at task3.out  
-----
```

- Run the program:

```
bash run.sh ./task3.out 2
```

```
-----  
Command executed: mpirun -np 2 ./task3.out  
-----  
#####  
#####          OUTPUT          #####  
#####  
#####  
  
Process 0 sent data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20  
Process 1 received data: 1 2 3 4 7 8 9 10 13 14 15 16  
  
#####  
#####          DONE          #####  
#####
```

## 15. task4 (transfer 3rd column to process 1)

```
#include <mpi.h>  
#include <stdio.h>  
#include <stdlib.h>  
#define N 8  
  
int main(int argc, char** argv) {  
    MPI_Init(&argc, &argv);  
    int rank, size;  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
    MPI_Comm_size(MPI_COMM_WORLD, &size);  
  
    if (size < 2) {  
        fprintf(stderr, "World size must be greater than 1 for this example\n");  
        MPI_Abort(MPI_COMM_WORLD, 1);  
    }  
    const int count = N;          // Number of blocks  
    const int blocklength = 1;    // Number of elements in each block  
    const int stride = N;         // Number of elements between the start of each block  
    MPI_Datatype vector_type;  
    MPI_Type_vector(count, blocklength, stride, MPI_INT, &vector_type);  
    MPI_Type_commit(&vector_type);  
    int data[N][N];               // Array to send/receive  
    if (rank == 0) {  
        for (int i = 0; i < N; i++) {
```

```

        for(int j = 0; j < N; j++){
            data[i][j] = i + 1;
        }
    }
    MPI_Send(&data[0][2], 1, vector_type, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent data: \n");
    for (int i = 0; i < N; i++) {
        for(int j = 0; j < N; j++){
            printf("%d ", data[i][j]);
        }
        printf("\n");
    }
    printf("\n");
} else if (rank == 1) {
    for (int i = 0; i < N; i++) {
        for(int j = 0; j < N; j++){
            data[i][j] = 0;
        }
    }
    MPI_Recv(&data[0][2], 1, vector_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received data: \n");
    for (int i = 0; i < N; i++) {
        for(int j = 0; j < N; j++){
            printf("%d ", data[i][j]);
        }
        printf("\n");
    }
    printf("\n");
}
MPI_Type_free(&vector_type);
MPI_Finalize();
return 0;
}

```

## 15.1. Compilation and Execution

- Compile the program:

```
bash compile.sh task4.c
```

```
-----
Command executed: mpicc task4.c -o task4.out
-----
```

```
Compilation successful. Check at task4.out
-----
```

- Run the program:

```
bash run.sh ./task4.out 2
```

```
-----  
Command executed: mpirun -np 2 ./task4.out  
-----
```

```
#####  
#####          OUTPUT          #####  
#####
```

```
Process 0 sent data:
```

```
1 1 1 1 1 1 1  
2 2 2 2 2 2 2  
3 3 3 3 3 3 3  
4 4 4 4 4 4 4  
5 5 5 5 5 5 5  
6 6 6 6 6 6 6  
7 7 7 7 7 7 7  
8 8 8 8 8 8 8
```

```
Process 1 received data:
```

```
0 0 1 0 0 0 0  
0 0 2 0 0 0 0  
0 0 3 0 0 0 0  
0 0 4 0 0 0 0  
0 0 5 0 0 0 0  
0 0 6 0 0 0 0  
0 0 7 0 0 0 0  
0 0 8 0 0 0 0
```

```
#####  
#####          DONE          #####  
#####
```

## 16. Implement vector using indexed

```
#include <mpi.h>  
#include <stdio.h>  
#include <stdlib.h>  
#define N 20  
  
int main(int argc, char** argv) {  
    MPI_Init(&argc, &argv);  
  
    int rank, size;  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
    MPI_Comm_size(MPI_COMM_WORLD, &size);  
  
    if (size < 2) {
```

```

    fprintf(stderr, "World size must be greater than 1 for this example\n");
    MPI_Abort(MPI_COMM_WORLD, 1);
}

const int count = 4;
const int sl = 5;           //sl = stridelenh
int blocklengths[4] = {2, 2, 2, 2};
int displacements[4] = {0 * sl, 1 * sl, 2 * sl, 3 * sl};
//int displacements[3] = {0, 5, 10};
int data[N];
MPI_Datatype indexed_type;

// Create an indexed datatype
MPI_Type_indexed(count, blocklengths, displacements, MPI_INT, &indexed_type);
MPI_Type_commit(&indexed_type);

if (rank == 0) {
    // Initialize the data array with some values
    for (int i = 0; i < N; i++) {
        data[i] = i + 1;
    }

    MPI_Send(data, 1, indexed_type, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent data: ");
    for (int i = 0; i < N; i++) {
        printf("%d ", data[i]);
    }
    printf("\n");
} else if (rank == 1) {
    // Initialize the data array to zero
    for (int i = 0; i < N; i++) {
        data[i] = 0;
    }

    MPI_Recv(data, 1, indexed_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received data: ");
    for (int i = 0; i < N; i++) {
        printf("%d ", data[i]);
    }
    printf("\n");
}

MPI_Type_free(&indexed_type);
MPI_Finalize();
return 0;
}

```

## 16.1. Compilation and Execution

- Compile the program:



```
bash compile.sh task5.c
```

```
-----  
Command executed: mpicc task5.c -o task5.out  
-----
```

```
Compilation successful. Check at task5.out  
-----
```

- Run the program:

```
bash run.sh ./task5.out 2
```

```
-----  
Command executed: mpirun -np 2 ./task5.out  
-----
```

```
#####  
#####          OUTPUT          #####  
#####
```

```
Process 0 sent data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
```

```
Process 1 received data: 1 2 0 0 0 6 7 0 0 0 11 12 0 0 0 16 17 0 0 0
```

```
#####  
#####          DONE          #####  
#####
```

## 17. Implement contiguous using indexed

```
#include <mpi.h>  
#include <stdio.h>  
#include <stdlib.h>  
#define N 20  
  
int main(int argc, char** argv) {  
    MPI_Init(&argc, &argv);  
  
    int rank, size;  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
    MPI_Comm_size(MPI_COMM_WORLD, &size);  
  
    if (size < 2) {  
        fprintf(stderr, "World size must be greater than 1 for this example\n");  
        MPI_Abort(MPI_COMM_WORLD, 1);  
    }  
}
```

```

const int count = 4;
const int sl = 5;           //sl = stridlength
int blocklengths[4] = {sl, sl, sl, sl}; //need to have stride length equal to block length to make it contiguous
int displacements[4] = {0 * sl, 1 * sl, 2 * sl, 3 * sl};
//int displacements[3] = {0, 5, 10};
int data[N];
MPI_Datatype indexed_type;

// Create an indexed datatype
MPI_Type_indexed(count, blocklengths, displacements, MPI_INT, &indexed_type);
MPI_Type_commit(&indexed_type);

if (rank == 0) {
    // Initialize the data array with some values
    for (int i = 0; i < N; i++) {
        data[i] = i + 1;
    }

    MPI_Send(data, 1, indexed_type, 1, 0, MPI_COMM_WORLD);
    printf("Process 0 sent data: ");
    for (int i = 0; i < N; i++) {
        printf("%d ", data[i]);
    }
    printf("\n");
} else if (rank == 1) {
    // Initialize the data array to zero
    for (int i = 0; i < N; i++) {
        data[i] = 0;
    }

    MPI_Recv(data, 1, indexed_type, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    printf("Process 1 received data: ");
    for (int i = 0; i < N; i++) {
        printf("%d ", data[i]);
    }
    printf("\n");
}

MPI_Type_free(&indexed_type);
MPI_Finalize();
return 0;
}

```

## 17.1. Compilation and Execution

- Compile the program:

```
bash compile.sh task6.c
```

```
-----
```

```
Command executed: mpicc task6.c -o task6.out
```

```
-----  
Compilation successful. Check at task6.out  
-----
```

- Run the program:

```
bash run.sh ./task6.out 2
```

```
-----  
Command executed: mpirun -np 2 ./task6.out  
-----
```

```
#####  
#####          OUTPUT          #####  
#####
```

```
Process 0 sent data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20  
Process 1 received data: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
```

```
#####  
#####          DONE          #####  
#####
```

Author: Abhishek Raj

Created: 2025-01-13 Mon 17:23