

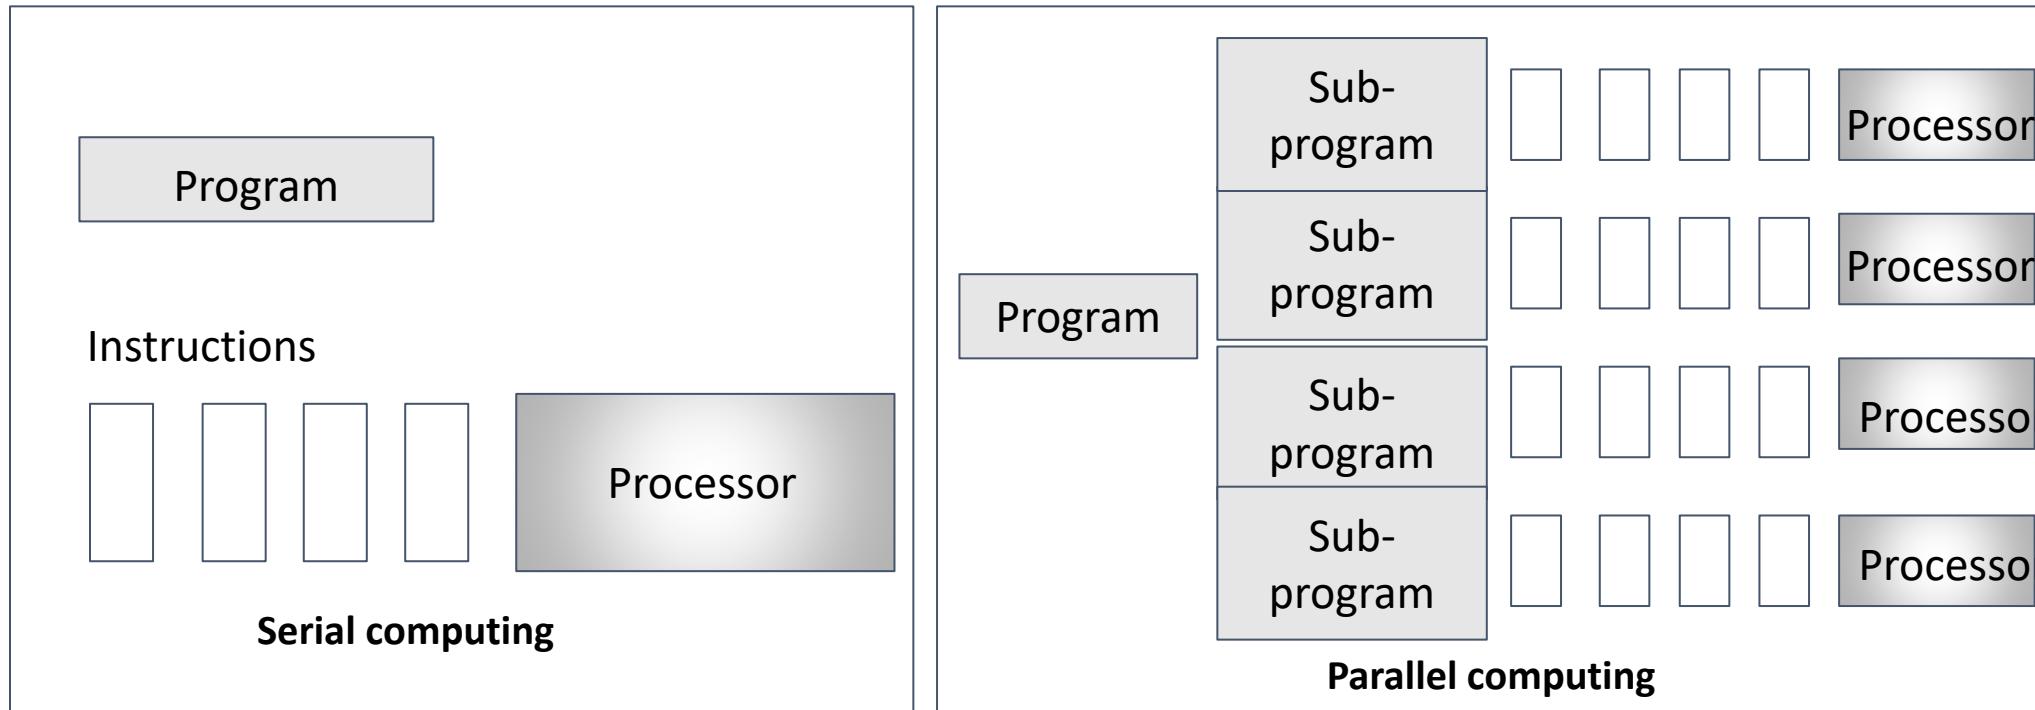
# Introduction to Parallel Programming

# Content

- Introduction to parallel programming
  - What is Parallel Programming?
  - Need of parallel programming.
  - Why parallel programming?
- Introduction to parallel hardware.
- General Parallel Computing Terminology
- Process and Threads.
- Demo

# Introduction to parallel programming

What is parallel programming?



# Introduction to parallel programming

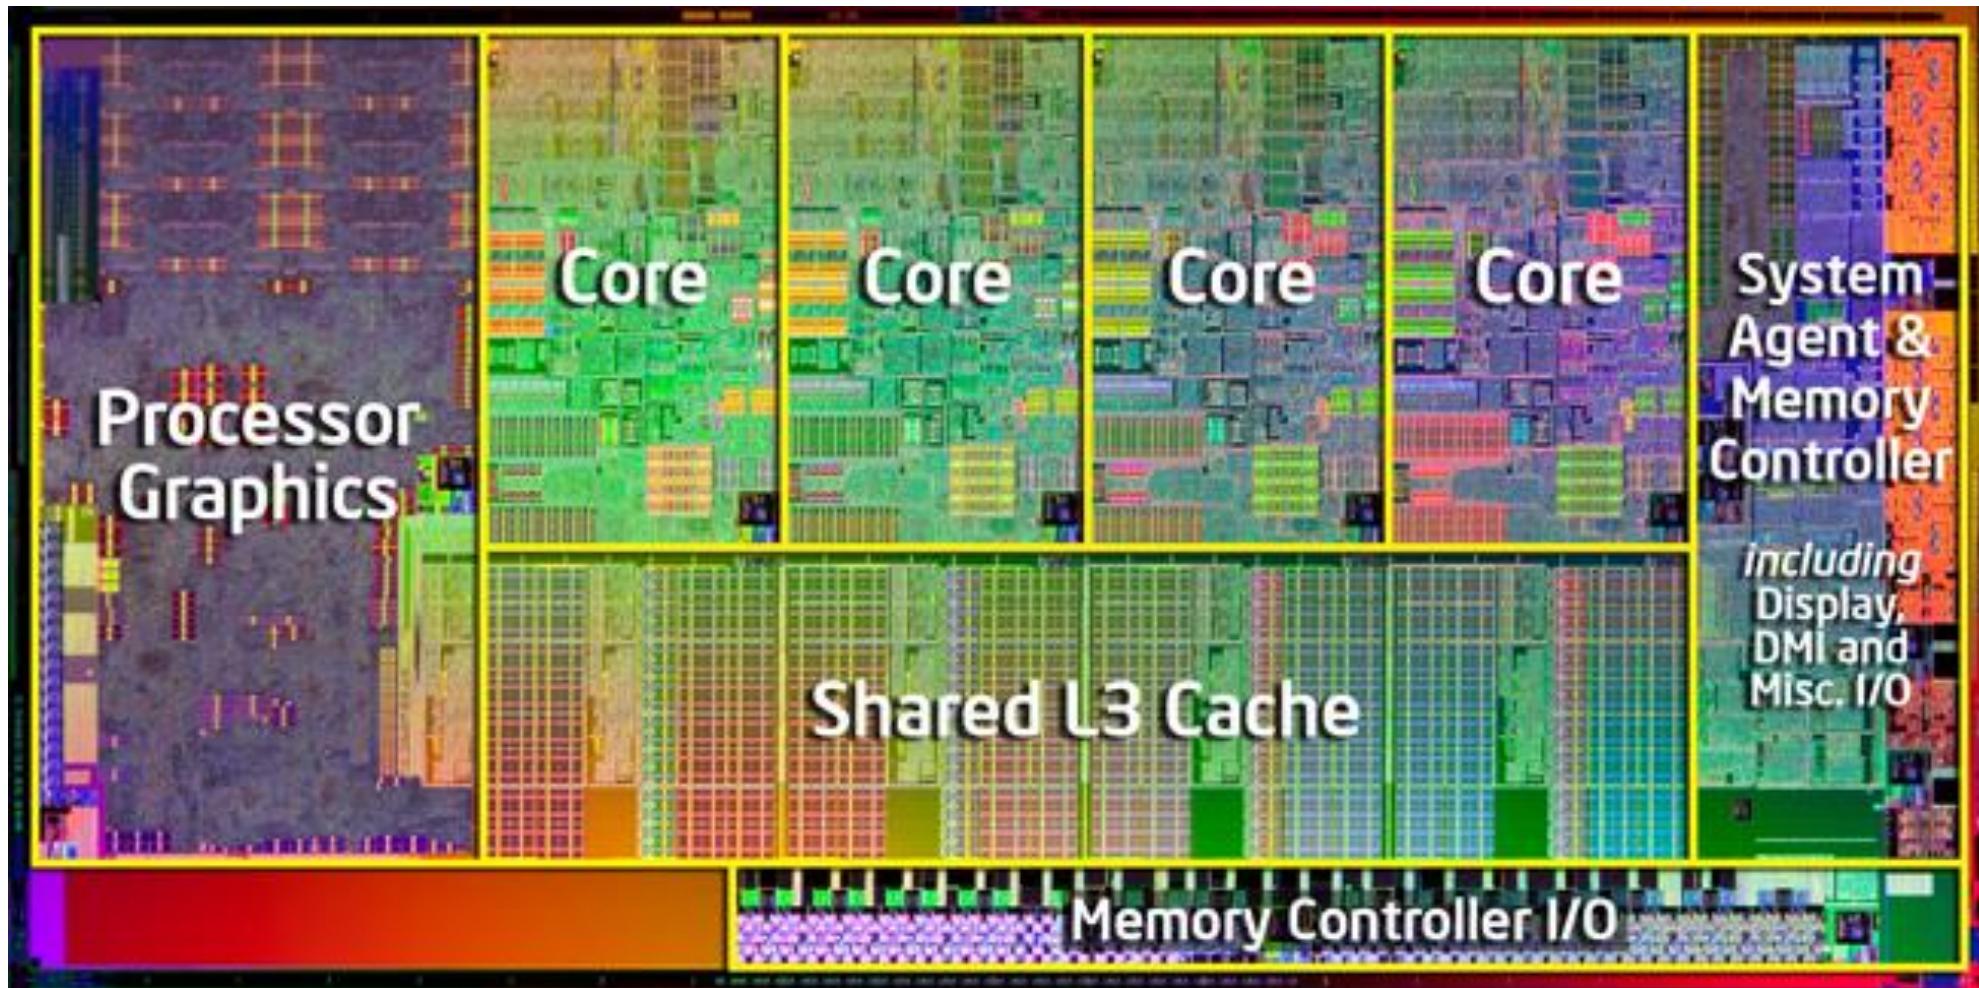
- Need of parallel programming.
  - The exponential growth of processing and network speeds means that parallel architecture isn't just a good idea; it's necessary. Big data and the IoT will soon force us to crunch trillions of data points at once.
  - Dual-core, quad-core, 8-core, and even 56-core chips are all examples of parallel computing. So, while parallel computers aren't new, here's the rub: new technologies are cranking out ever-faster networks, and computer performance has grown 250,000 times in 20 years.
  - For instance, in just the healthcare sector, AI tools will be rifling through the heart rates of a hundred million patients, looking for the telltale signs of A-fib or V-tach and saving lives. They won't be able to make it work if they have to plod along performing one operation at a time

# Introduction to parallel programming

Why parallel programming?

- PARALLEL COMPUTING MODELS THE REAL WORLD.
- SAVES TIME
- SAVES MONEY
- SOLVE MORE COMPLEX OR LARGER PROBLEMS
- LEVERAGE REMOTE RESOURCES

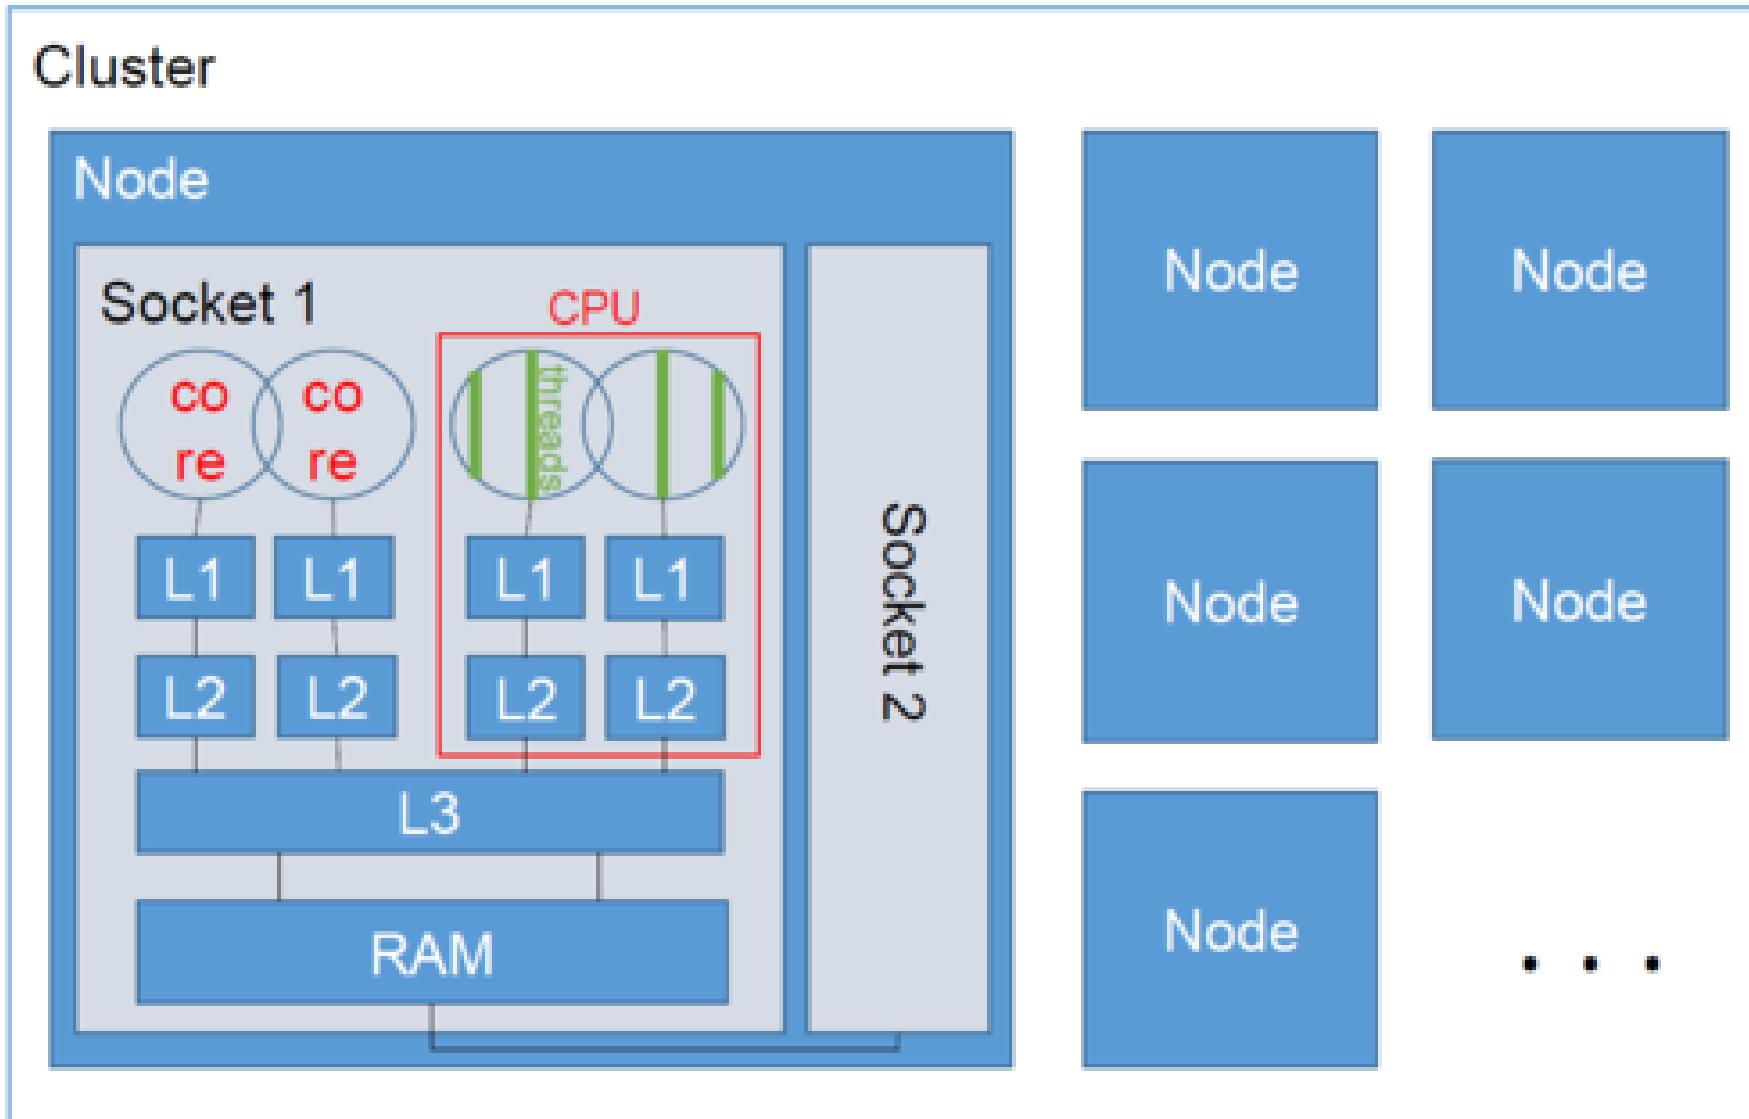
# Introduction to parallel hardware.



# General Parallel Computing Terminology

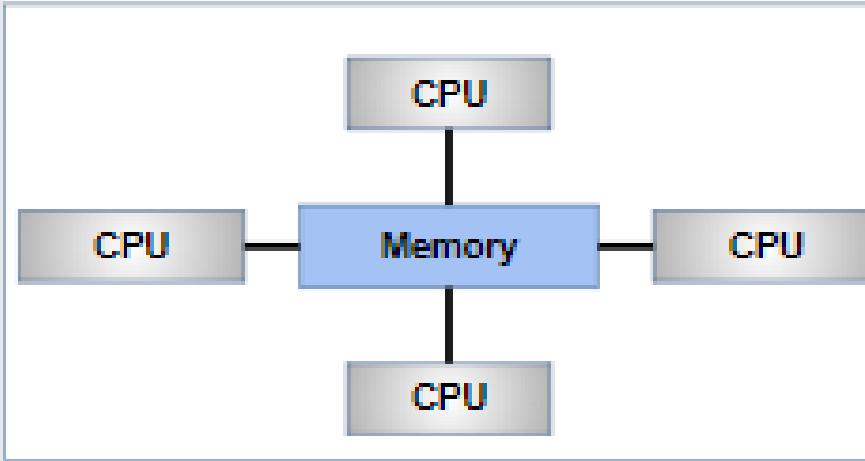
- Node
  - A standalone "computer in a box." Usually comprised of multiple CPUs/processors/cores, memory, network interfaces, etc. Nodes are networked together to comprise a supercomputer.
- CPU
  - Contemporary CPUs consist of one or more cores.
- Process
  - A process is an instance of a program that is being executed or processed
- Thread
  - Thread is a segment of a process or a lightweight process that is managed by the scheduler independently

# General Parallel Computing Terminology

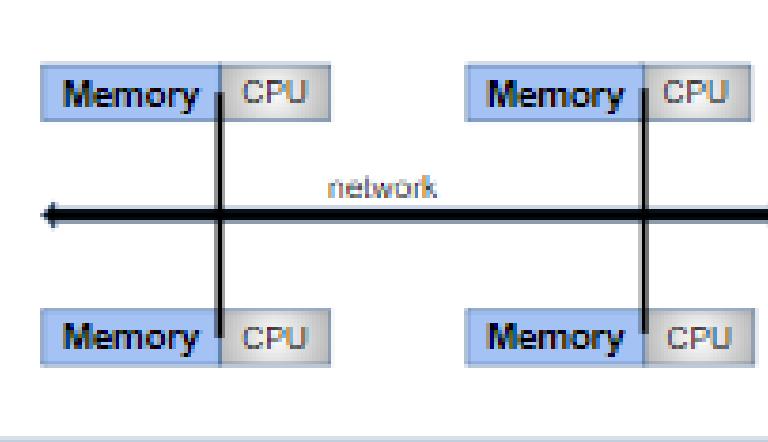


# General Parallel Computing Terminology

**Shared Memory**



**Distributed Memory**

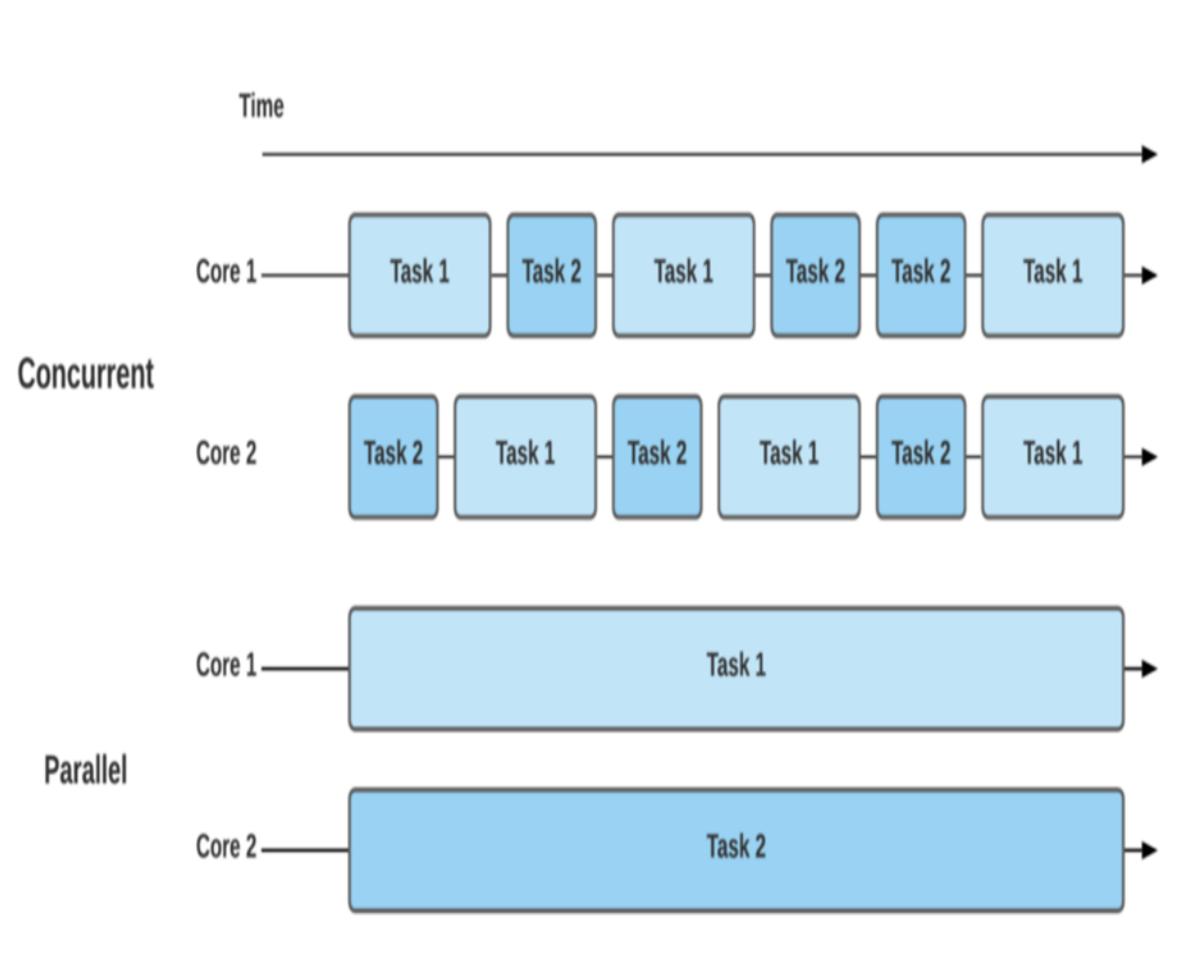


Openmp

Message passing  
interface

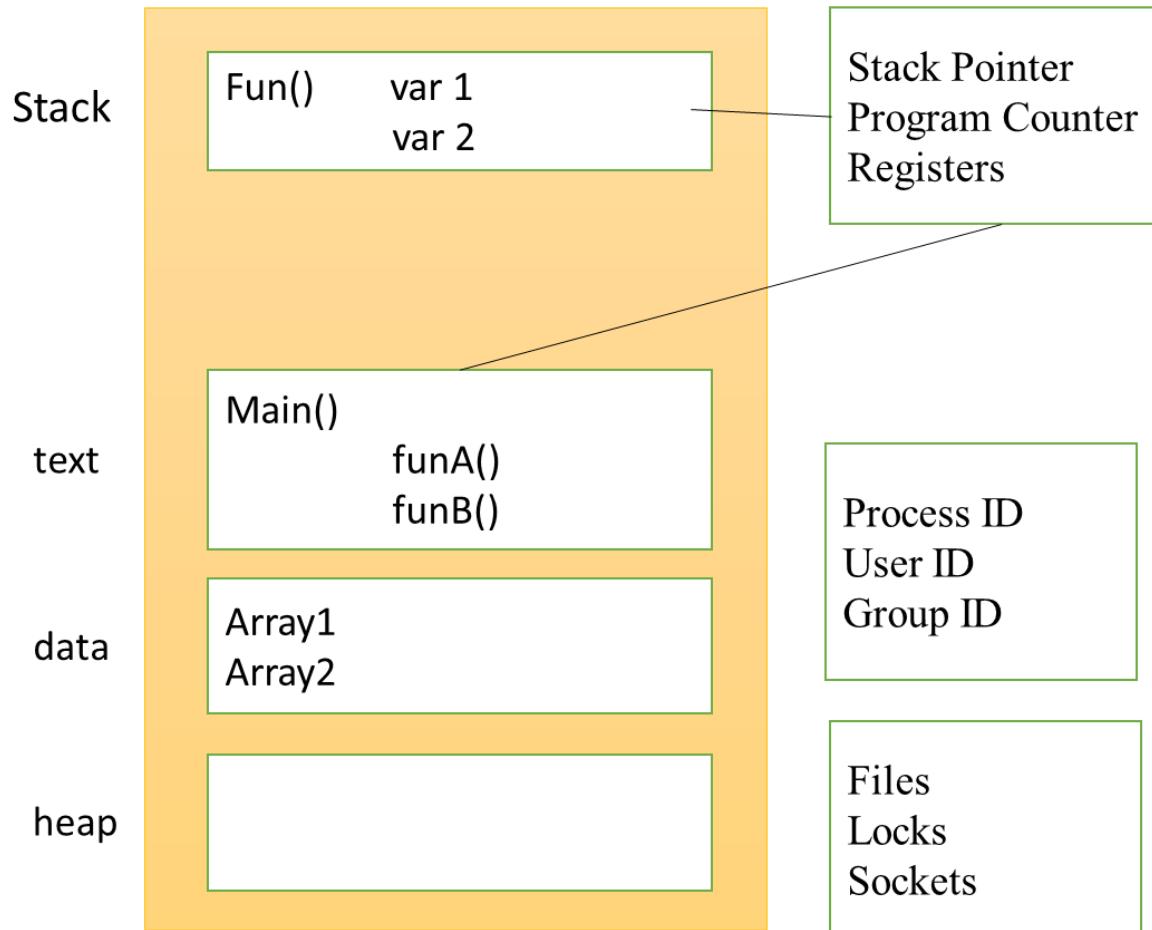
# General Parallel Computing Terminology

- Concurrency and Parallelism
  - Concurrency: A condition of a system in which multiple tasks are logically active at one time
  - Parallelism: A condition of a system in which multiple tasks are actually active at one time



# Process and Threads

- How program execute in memory?

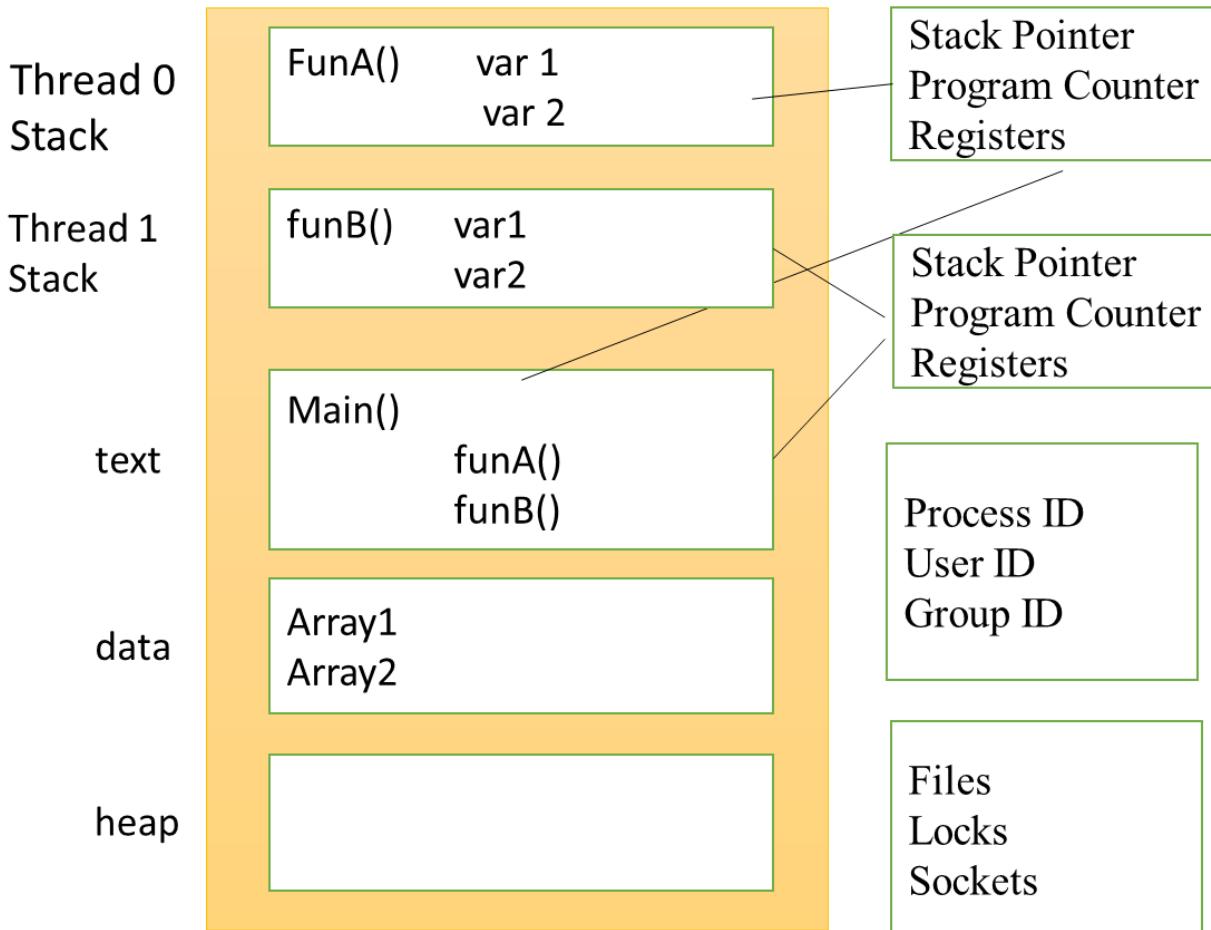


## Process:

- An instance of a program execution.
- The execution context of a running program ... i.e. the resources associated with a program's execution.

# Process and Threads

- How program execute in shared memory?



## Threads:

- Threads are "light weight processes"
- Threads share Process state among multiple threads ... this greatly reduces the cost of switching context.

# Demo Time



# Hardware information

\$lscpu or cat /proc/cpuinfo

```
(base) [cdacapp@login03 ~]$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:             Little Endian
CPU(s):                40
On-line CPU(s) list:   0-39
Thread(s) per core:    1
Core(s) per socket:    20
Socket(s):              2
NUMA node(s):           2
Vendor ID:              GenuineIntel
CPU family:             6
Model:                 85
Model name:             Intel(R) Xeon(R) Gold 6248 CPU @ 2.50GHz
Stepping:               7
CPU MHz:                999.908
CPU max MHz:            3900.0000
CPU min MHz:            1000.0000
BogoMIPS:               5000.00
Virtualization:         VT-x
L1d cache:              32K
L1i cache:              32K
L2 cache:                1024K
L3 cache:                28160K
NUMA node0 CPU(s):      0-19
NUMA node1 CPU(s):      20-39
Flags:                  fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts ac
pi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc art arch_perfmon pebs bts rep_goo
d noopl xtopology nonstop_tsc aperfmpfperf eagerfpu pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 sd
bg fma cx16 xtpr pdcm pcid dca sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand
lahf_lm abm 3dnowprefetch epb cat_l3 cdp_l3 invpcid_single intel_ppin intel_pt ssbd mba ibrs ibpb stibp ibrs
_enhanced tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm
cqm mpx rdt_a avx512f avx512dq rdseed adx smap clflushopt clwb avx512cd avx512bw avx512vl xsaveopt xsavec xge
tbv1 cqm_llc cqm_occup_llc cqm_mbm_total cqm_mbm_local dtherm ida arat pln pts pku ospke avx512_vnni md_clear
spec_ctrl intel_stibp flush l1d arch_capabilities
(base) [cdacapp@login03 ~]$ █
```

# Process and Thread information

\$top -u <username>

| PID   | USER    | PR | NI | VIRT    | RES    | SHR  | S | %CPU | %MEM | TIME+   | COMMAND |
|-------|---------|----|----|---------|--------|------|---|------|------|---------|---------|
| 43243 | cdacapp | 20 | 0  | 167096  | 3204   | 1664 | R | 1.0  | 0.0  | 0:00.09 | top     |
| 28610 | cdacapp | 20 | 0  | 6498508 | 296596 | 1560 | S | 0.0  | 0.1  | 0:09.44 | julia   |
| 28611 | cdacapp | 20 | 0  | 6493920 | 300736 | 1564 | S | 0.0  | 0.1  | 0:08.59 | julia   |
| 28612 | cdacapp | 20 | 0  | 6569328 | 327268 | 1592 | S | 0.0  | 0.1  | 0:08.79 | julia   |
| 28613 | cdacapp | 20 | 0  | 6565220 | 330992 | 1568 | S | 0.0  | 0.1  | 0:09.04 | julia   |
| 42747 | cdacapp | 20 | 0  | 167760  | 2724   | 1208 | S | 0.0  | 0.0  | 0:00.04 | sshd    |
| 42748 | cdacapp | 20 | 0  | 121768  | 4188   | 1728 | S | 0.0  | 0.0  | 0:00.18 | bash    |

# Program in C

```
#include <stdio.h>
int main() {
    int n, i, sum = 0;

    printf("Enter a positive integer: ");
    scanf("%d", &n);

    for (i = 1; i <= n; ++i) {
        sum += i;
    }

    printf("Sum = %d", sum);
    return 0;
}
```

How to compile:

\$gcc <file\_name>.c –o <file\_name>  
\$time <file\_name>

Thank you

# Introduction to OPENMP

Open Multi Processing

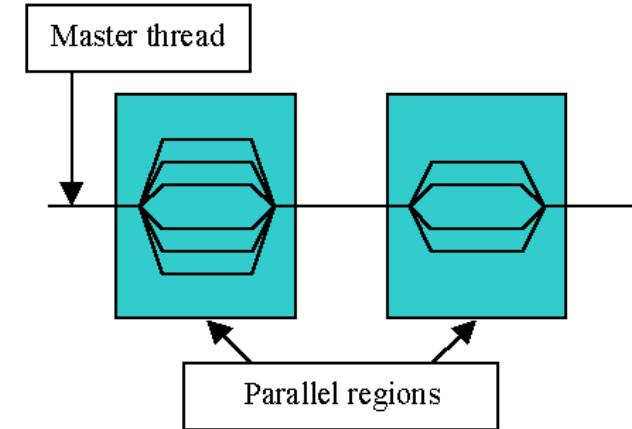
# Content

- Introduction to openmp
  - What is OpenMP?
  - History of OpenMP
  - Why openmp
- OpenMP Programming Model
- OpenMP Stack
- Hello world in openmp
  - Basic Syntax
  - Hello world c program
  - Compile and execution

# Introduction to openmp

## What is OpenMP?

- An Application Program Interface (API) that may be used to explicitly direct multi-threaded, shared memory parallelism
- Comprises three primary API components
  - Compiler Directives
  - Runtime Library Routines
  - Environment Variables
- Portable
  - The API is specified for C/C++ and Fortran
  - Has been implemented for most major platforms including Unix/ Linux platforms and Windows NT



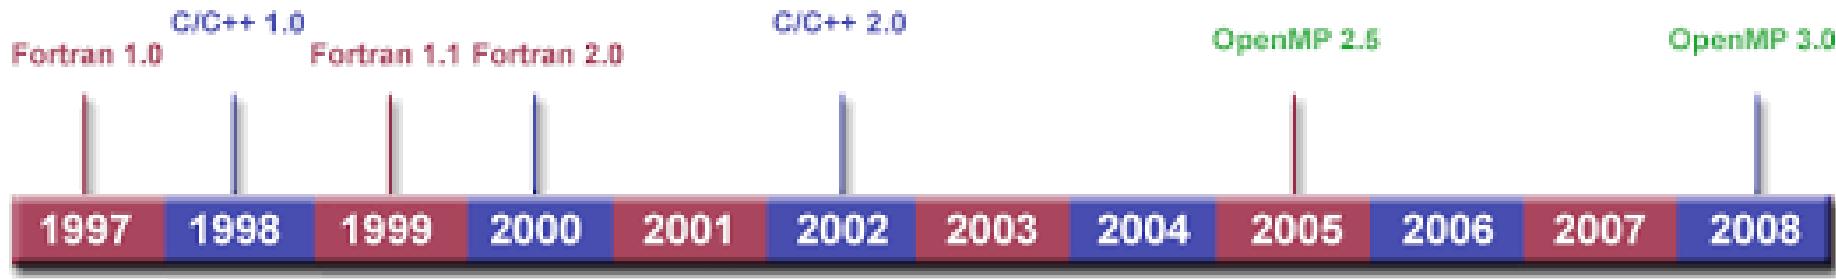
# Introduction to openmp

## What is OpenMP? (cont.)

- Standardized
  - Jointly defined and endorsed by a group of major computer hardware and software vendors
  - Expected to become an ANSI standard later???
- What does OpenMP stand for?
  - Short version: Open Multi-Processing
  - Long version: Open specifications for Multi-Processing via collaborative work between interested parties from the hardware and software industry, government and academia

# Introduction to openmp

## History of OpenMP



Version 1.0 : OpenMP for Fortran 1.0, in October 1997

Version 1.0 : OpenMP for C/C++, in October 1998

Version 2.0 : OpenMP for Fortran 1.1, in 1999

Version 2.0 : OpenMP for C/C++, in 2000

Version 2.0 : OpenMP for C/C++, in 2002

Version 2.5 : OpenMP for C/C++/Fortran, in 2005

# Introduction to openmp

## History of OpenMP(cont.)

- Up to version 2.0, OpenMP primarily specified ways to parallelize highly regular loops, as they occur in matrix-oriented numerical programming,
- In Version 3.0 Included features like concept of tasks and the task construct.
- In version 4.0 openmp improved following features:
  - support for accelerators, atomics, error handling, thread affinity; tasking extensions; user defined reduction, SIMD support.



# Introduction to openmp

## Why Openmp?

- More efficient
- Hides the low-level details.
- OpenMP has directives that allow the programmer to:
  - specify the parallel region
  - specify whether the variables in the parallel section are private or shared
  - specify how/if the threads are synchronized
  - specify how to parallelize loops
  - specify how the works is divided between threads (scheduling)

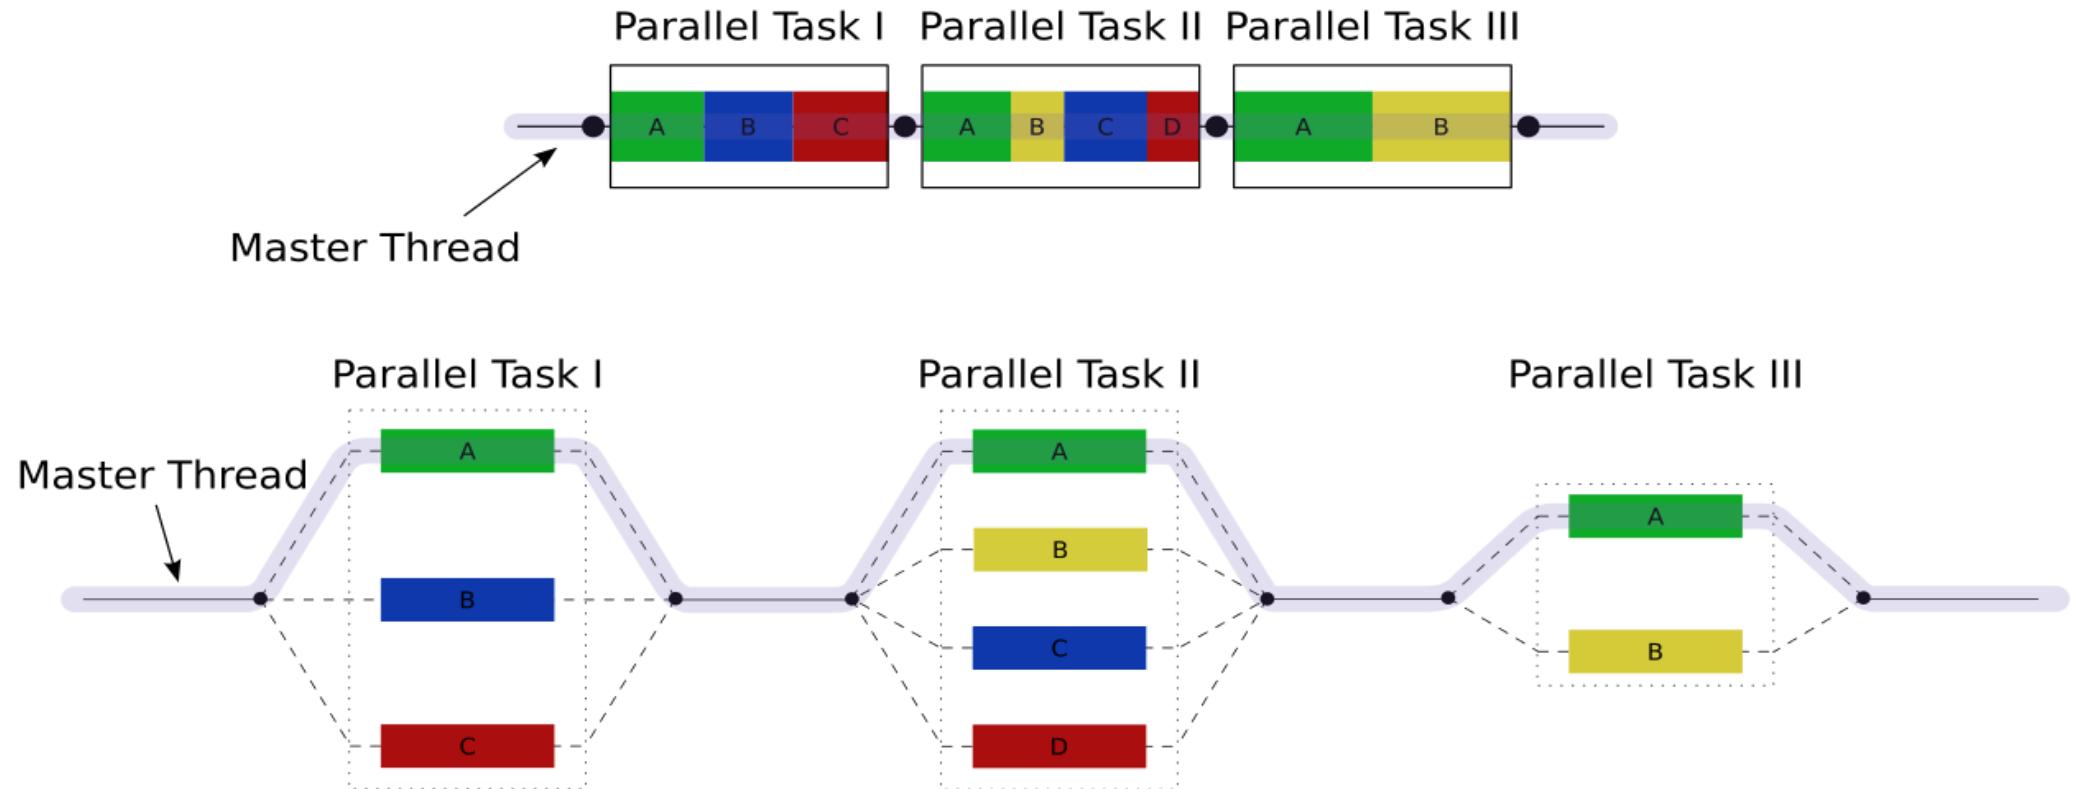


# OpenMP Programming Model

- Shared memory, thread-based parallelism
  - OpenMP is based upon the existence of multiple threads in the shared memory programming paradigm.
  - A shared memory process consists of multiple threads.
- Explicit Parallelism
  - OpenMP is an explicit (not automatic) programming model, offering the programmer full control over parallelization.
- OpenMP uses the fork-join model of parallel execution.
- Compiler directive based
  - Most OpenMP parallelism is specified through the use of compiler directives which are imbedded in C/C++ or Fortran source code.

# OpenMP Programming Model

## Fork–join model:



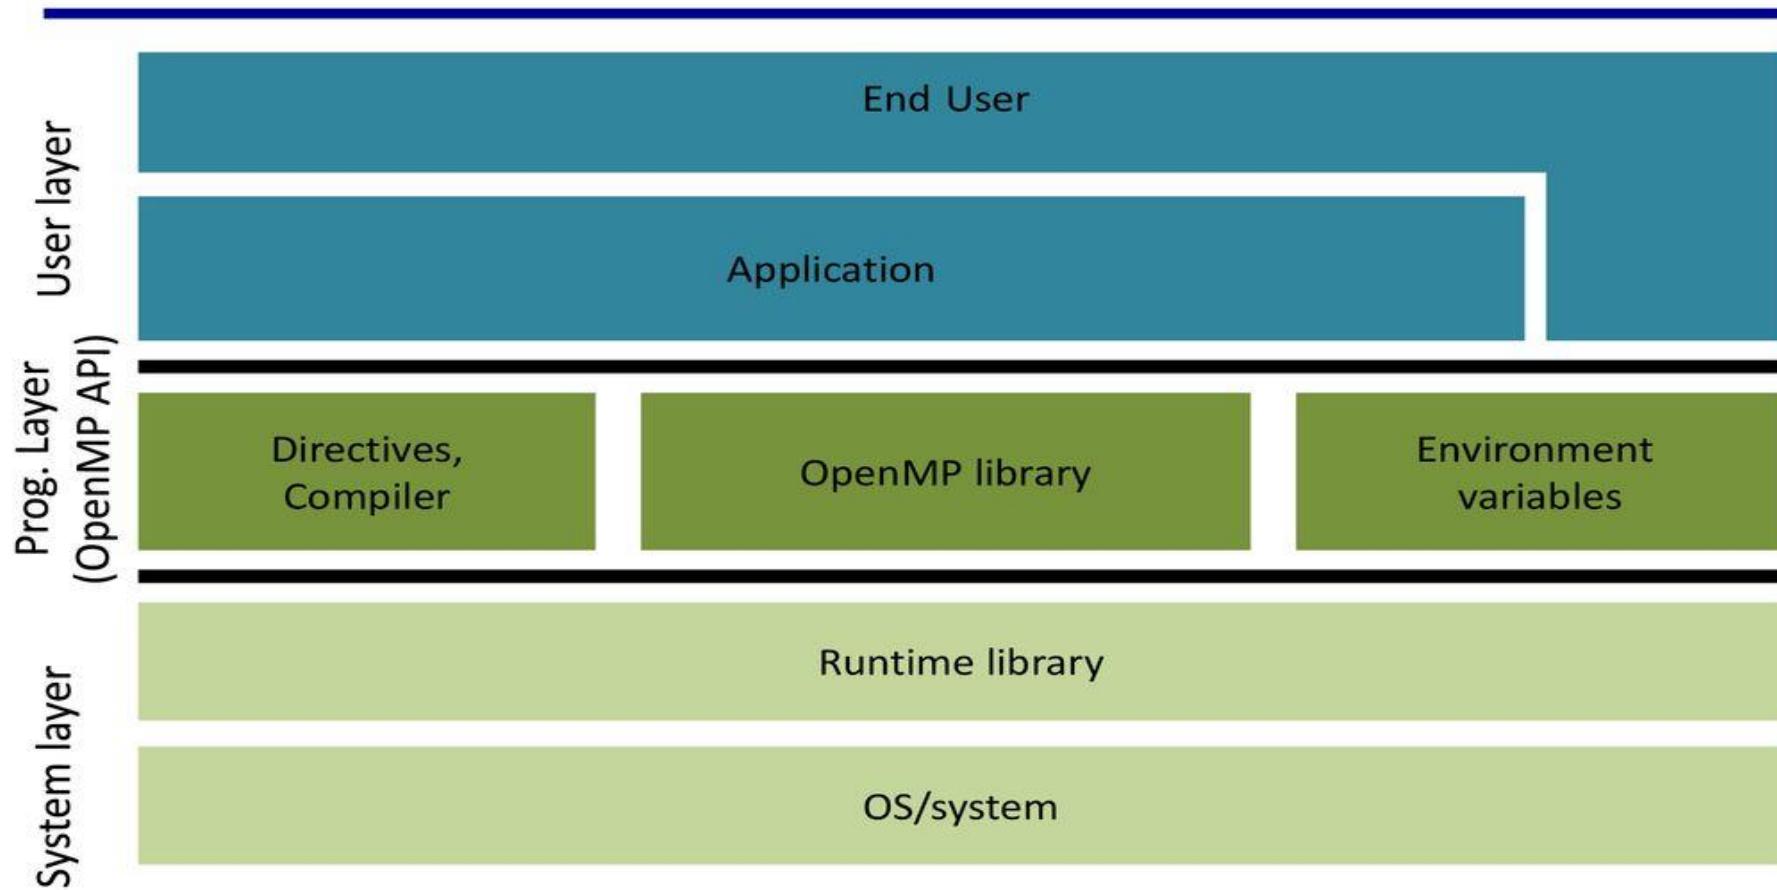
# OpenMP Programming Model

## Fork–join model: (cont.)

- All OpenMP programs begin as a single process: the master thread. The **master thread** executes sequentially until the first parallel region construct is encountered.
- **FORK**: the master thread then creates a **team** of parallel threads
- The statements in the program that are enclosed by the parallel region construct are then executed in parallel among the team threads.
- **JOIN**: When the team threads complete the statements in the parallel region construct, they synchronize and terminate, leaving only the master thread

# OpenMP Stack

## OpenMP Parallel Computing Solution Stack



# Hello world in openmp

## Basic Syntax

- Function prototypes and types in the file:

```
#include <omp.h>
```

- Most of the constructs in OpenMP are compiler directives.

```
#pragma omp construct [clause [clause]...]  
{  
    //..Do some work here  
}  
//end of parallel region/block
```

- Example:

```
#pragma omp parallel num_threads(4)
```

# Hello world in openmp

## OpenMP Hello World program using C

The diagram illustrates an OpenMP Hello World program in C. The code is shown in a dark background with numbered lines. Three callout boxes provide explanations for specific parts of the code:

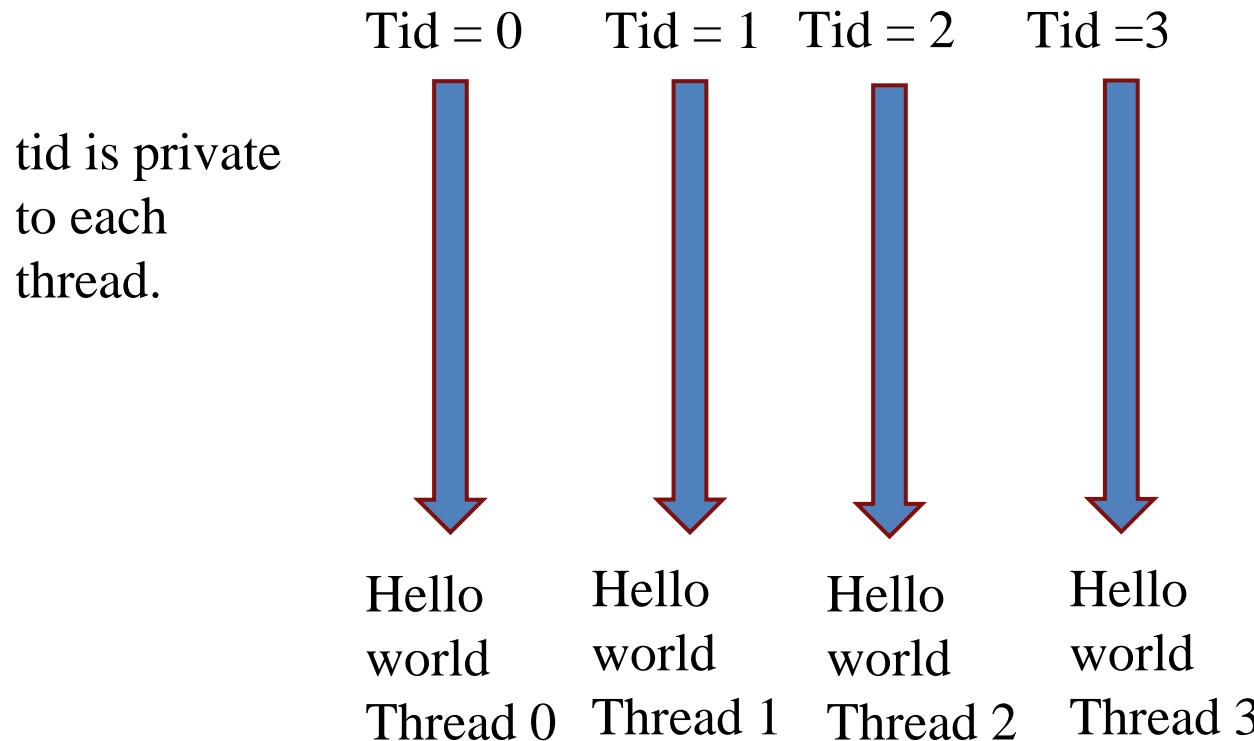
- OpenMP Include File**: Points to the line `#include<omp.h>`. This is highlighted in red.
- Runtime library function to return a thread ID.**: Points to the line `int ID = omp_get_thread_num();`. This is highlighted in green.
- Parallel region with default number of threads**: Points to the line `#pragma omp parallel {`. This is highlighted in yellow.

```
1 #include<stdio.h>
2 #include<omp.h>
3
4 int main(void)
5 {
6
7     #pragma omp parallel
8     {
9         int ID = omp_get_thread_num();
10        printf("Hello, world(%d)\n", ID);
11    }
12
13    return 0;
14
15 }
```

# Hello world in openmp

**Number of threads 4**

Parallel Region



# Compile and execution

## Compile program in OpenMP

- Set number of threads:

Using shell

```
$ export OMP_NUM_THREADS=4
```

Inside program before parallel region.

```
omp_set_num_threads(4);
```

- For GNU C compiler:

```
$ gcc -fopenmp HelloWorld.c -o Hello
```

```
./Hello
```

- For Intel compiler:

```
$ icc -qopenmp HelloWorld.c -o Hello
```

```
./Hello
```

## Output of hello world program with 4 threads:

```
[parikshita@shavak solutions]$ ./Hello
Hello World(3)
Hello World(0)
Hello World(1)
Hello World(2)
[parikshita@shavak solutions]$
```



**THANK YOU**

**FOR YOUR ATTENTION**

# OpenMP Components

## Language extensions

# Content

## OpenMP Components

- PARALLEL Region Construct.
  - Number of Threads
  - Nested Parallel Regions.
  - PARALLEL Region Example
  - Nested Parallel Region Example
- Work-sharing Constructs
  - Do/for
  - Single
  - Section
- Synchronization

# OpenMP Components

## Directives

- Parallel region
- Worksharing constructs
- Tasking
- Offloading
- Affinity
- Error Handling
- SIMD
- Synchronization
- Data-Sharing attributes

## Runtime Environment

- Number of threads
- Thread Id
- Dynamic thread adjustment
- Nested parallelism
- Schedule
- Thread limit
- Wallclock timer
- Team size
- Nesting level

## Runtime Variables

- Number of threads
- Scheduling type
- Dynamic thread adjusting
- Nested parallelism
- Thread limit

# Parallel Region Construct

- Block of code that will be executed by multiple threads
- Fundamental OpenMP parallel construct
- When a thread reaches a PARALLEL directive, it creates a team of threads and becomes the master of the team. The master is a member of that team and has thread number 0 within that team.
- Starting from the beginning of this parallel region, the code is duplicated and all threads will execute that code.
- There is an implied barrier at the end of a parallel section Only the master thread continues execution past this point.
- If any thread terminates within a parallel region, all threads in the team will terminate, and the work done up until that point is undefined.

# Parallel Region Construct

**#pragma omp parallel [clause[ [, ]clause] ...] new-line**

clause:

if(scalar-expression)

num\_threads(integer-expression)

default(shared | none)

private(list)

firstprivate(list)

shared(list)

reduction(operator: list)

# Parallel Region Construct

## Number of Threads

- The number of threads in a parallel region is determined by the following factors, in order of precedence:
  1. Evaluation of the IF clause
  2. Setting of the NUM\_THREADS clause
  3. Use of the `omp_set_num_threads()` library function
  4. Setting of the OMP\_NUM\_THREADS environment variable
  5. Implementation default - usually the number of cores on a node.
- Threads are numbered from 0 (master thread) to N-1

# PARALLEL Region Example

```
#include <omp.h>
#include <stdio.h>
int main (){
#pragma omp parallel
if(omp_in_parallel)
{
    printf("The threads is
%d\n",omp_get_thread_num());
}
return 0;
}
```

How to compile and run?  
\$gcc -fopenmp <file\_name.c> -o  
<file\_name>  
\$./<file\_name>

```
[cdacapp2@login3 1.lf]$ ./mp_if
The threads is 0
The threads is 2
The threads is 3
The threads is 1
```

# Parallel Region Construct

## Nested Parallel Regions

- Use the `omp_get_nested()` library function to determine if nested parallel regions are enabled.
- The two methods available for enabling nested parallel regions (if supported) are:
  - The `omp_set_nested()` library routine
  - Setting of the `OMP_NESTED` environment variable to `TRUE`
- If not supported, a parallel region nested within another parallel region results in the creation of a new team, consisting of one thread, by default.

# PARALLEL Nested Region Example

```
#include <omp.h>
#include <stdio.h>
void report_num_threads(int level)
{
    #pragma omp single
    {
        printf("Level %d: number of threads in the team - %d\n",
               level, omp_get_num_threads());
    }
}
int main()
{
    omp_set_dynamic(0);
    #pragma omp parallel num_threads(2)
    {
        report_num_threads(1);
        #pragma omp parallel num_threads(2)
        {
            report_num_threads(2);
            #pragma omp parallel num_threads(2)
            {
                report_num_threads(3);
            }
        }
    }
    return(0);
}
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o
<file_name>
$./<file_name>
```

```
(base) [cdacapp@login01 nested]$ export OMP_NESTED=True
(base) [cdacapp@login01 nested]$ ./nested
Level 1: number of threads in the team - 2
Level 2: number of threads in the team - 2
Level 2: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
```

# Work-sharing Constructs

- A work-sharing construct divides the execution of the enclosed code region among the members of the team that encounter it.
- Work-sharing constructs do not launch new threads
- There is no implied barrier upon entry to a work-sharing construct, however there is an implied barrier at the end of a work sharing construct.
- A work-sharing construct must be enclosed dynamically within a parallel region in order for the directive to execute in parallel.
- Work-sharing constructs must be encountered by all members of a team or none at all
- Successive work-sharing constructs must be encountered in the same order by all members of a team.

# Types of Work-sharing Constructs

- DO / for
  - Shares iterations of a loop across the team
  - Represents a type of "data parallelism"
- SECTIONS
  - Breaks work into separate, discrete sections
  - Each section is executed by a thread.
  - Can be used to implement a type of "functional parallelism"
- SINGLE
  - Serializes a section of code

# For Directive Syntax

**#pragma omp for [clause[,] clause] ... ] new-line**

for-loops

clause:

private(list)

firstprivate(list)

lastprivate(list)

reduction(operator: list)

schedule(kind[, chunk\_size])

collapse(n)

ordered

nowait

# For loop

The for construct tells OpenMP that the iteration set of the for loop that follows is to be distributed across the threads present in the team. Without the for construct, the entire iteration set of the for loop concerned will be executed by each thread in the team.

**The for directive supports the following clauses:**

- **private**
- **firstprivate**
- **lastprivate**
- **reduction**
- **ordered**
- **schedule**
- **Nowait**

- How iteration are divide in threads?  
Example: consider we have 2 threads  
and N iterations then

Threads 0 - 0 -  $N/2-1$

Threads 1 -  $N/2$  -  $N-1$

## For loop - PRIVATE

- The values of private data are undefined upon entry to and exit from the specific construct.
- Loop iteration variable is private by default.

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$.</><file_name>
```

```
$ ./for_private  
The thread 0 value is 0  
The thread 0 value is 30  
The thread 1 value is 0  
The thread 1 value is 30  
The thread 3 value is 0  
The thread 3 value is 30  
The thread 4 value is 0  
The thread 4 value is 30  
The thread 2 value is 0  
The thread 2 value is 30
```

# For loop - PRIVATE



```
1 #include<stdio.h>
2 #include<omp.h>
3 #define N 5
4 int main()
5 {
6     int a = 10; //shared
7     int b = 20; //shared
8     int c = 20;
9     omp_set_num_threads(N);
10    #pragma omp parallel for private(c)
11    for(int i = 0; i < N;i++){
12        printf("The thred %d value is %d\n",omp_get_thread_num(),c);
13        c = a + b;      //private
14        printf("The thred %d value is %d\n",omp_get_thread_num(),c);
15    }
16    return 0;
17 }
```

## For loop - FIRSTPRIVATE

- The clause combines behavior of private clause with automatic initialization of the variables in its list.

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

```
$ ./for_firstprivate  
The thred 0 value is 20  
The thred 0 value is 30  
The thred 4 value is 20  
The thred 4 value is 30  
The thred 3 value is 20  
The thred 3 value is 30  
The thred 2 value is 20  
The thred 2 value is 30  
The thred 1 value is 20  
The thred 1 value is 30
```

# For loop - FIRSTPRIVATE



```
1 #include<stdio.h>
2 #include<omp.h>
3 #define N 5
4 int main()
5 {
6     int a = 10; //shared
7     int b = 20; //shared
8     int c = 20;
9     omp_set_num_threads(N);
10    #pragma omp parallel for firstprivate(c)
11    for(int i = 0; i < N;i++){
12        printf("The thred %d value is %d\n",omp_get_thread_num(),c);
13        c = a + b; //private
14        printf("The thred %d value is %d\n",omp_get_thread_num(),c);
15    }
16    return 0;
17 }
```

## For loop - LASTPRIVATE

- Performs finalization of private variables
- Each thread has its own copy

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

```
$ ./for_lastprivate  
The thred 4 value is 30  
The thred 2 value is 30  
The thred 1 value is 30  
The thred 3 value is 30  
The thred 0 value is 30  
The thred 0 value is 30
```

# For loop - LASTPRIVATE



```
1 #include<stdio.h>
2 #include<omp.h>
3 #define N 5
4 int main()
5 {
6     int a = 10; //shared
7     int b = 20; //shared
8     int c ;
9     omp_set_num_threads(N);
10    #pragma omp parallel for lastprivate(c)
11    for(int i = 0; i < N;i++){
12        c = a + b;      //private
13        printf("The thred %d value is %d\n",omp_get_thread_num(),c);
14    }
15    printf("The thred %d value is %d\n",omp_get_thread_num(),c);
16    return 0;
17 }
```

## For loop - REDUCTION

- The reduction clause indicates that the variables passed are, as its name suggests, used in a reduction.
- By default, the reduction computation is complete at the end of the construct.

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

\$ ./reduction

Number of primes numbers:

5761455

# For loop - REDUCTION

```
#include <stdio.h>
#include <omp.h>
#define NUM 100000000
int isprime( int x )
{
    for( int y = 2; y * y <= x; y++ )
    {
        if( x % y == 0 )
            return 0;
    }
    return 1;
}
int main( )
{
    int sum = 0;
#pragma omp parallel for reduction (+:sum)
    for( int i = 2; i <= NUM ; i++ )
    {
        sum += isprime ( i );
    }
    printf( "Number of primes numbers: %d\n", sum );
    return 0;
}
```

## For loop - ORDERED

- The `omp ordered` directive identifies a structured block of code that must be executed in sequential order.

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

```
$ ./ordered
```

```
Thread 1 processes iteration 3.  
Thread 1 processes iteration 4.  
Thread 1 processes iteration 5.  
Thread 2 processes iteration 6.  
Thread 2 processes iteration 7.  
Thread 0 processes iteration 0.  
Thread 0 processes iteration 1.  
Thread 0 processes iteration 2.  
Thread 3 processes iteration 8.  
Thread 3 processes iteration 9.
```

# For loop - ORDERED



```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <omp.h>
4
5 int main(int argc, char* argv[])
6 {
7     omp_set_num_threads(4);
8     #pragma omp parallel for ordered
9     for(int i = 0; i < 10; i++)
10    {
11        printf("Thread %d processes iteration %d.\n", omp_get_thread_num(), i);
12    }
13
14    return 0;
15 }
```

## For loop - SCHEDULE

- The schedule clause tells OpenMP how to distribute the loop iterations to the threads.

**The OpenMP scheduling kind to use.**

**Possible values:**

- auto: the auto scheduling kind will apply.
- dynamic: the dynamic scheduling kind will apply.
- guided: the guided scheduling kind will apply.
- runtime: the runtime scheduling kind will apply.
- static: the static scheduling kind will apply.

## SCHEDULE Clause.. cont

### **schedule(static, [n])**

- Each thread is assigned chunks in “round robin” fashion, known as block cyclic scheduling.
- divides iterations into chunks that are approximately equal in size and it distributes at most one chunk to each thread
- If n has not been specified, it will contain  $\text{CEILING}(\text{number\_of\_iterations} / \text{number\_of\_threads})$  iterations

Example: loop of length 16, 3 threads, chunk size 2 :

# For loop - SCHEDULE(static)

```
 1 #include <stdio.h>
 2 #include <stdlib.h>
 3 #include <omp.h>
 4
 5 int main(int argc, char* argv[])
 6 {
 7     omp_set_num_threads(3);
 8
 9     printf("With no chunksize passed:\n");
10     #pragma omp parallel for schedule(static)
11     for(int i = 0; i < 10; i++)
12     {
13         printf("Thread %d processes iteration %d.\n", omp_get_thread_num(), i);
14     }
15     printf("With a chunksize of 2:\n");
16
17     #pragma omp parallel for schedule(static, 2)
18     for(int i = 0; i < 10; i++)
19     {
20         printf("Thread %d processes iteration %d.\n", omp_get_thread_num(), i);
21     }
22
23     return 0;
24 }
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>
$./<file_name>
```

\$ ./parallel\_for

With no chunksize passed:

Thread 0 processes iteration 0.  
Thread 0 processes iteration 1.  
Thread 0 processes iteration 2.  
Thread 2 processes iteration 7.  
Thread 2 processes iteration 8.  
Thread 2 processes iteration 9.  
Thread 0 processes iteration 3.  
Thread 1 processes iteration 4.  
Thread 1 processes iteration 5.  
Thread 1 processes iteration 6.

With a chunksize of 2:

Thread 1 processes iteration 2.  
Thread 1 processes iteration 3.  
Thread 1 processes iteration 8.  
Thread 1 processes iteration 9.  
Thread 2 processes iteration 4.  
Thread 2 processes iteration 5.  
Thread 0 processes iteration 0.  
Thread 0 processes iteration 1.  
Thread 0 processes iteration 6.  
Thread 0 processes iteration 7.

# For loop - SCHEDULE(dynamic)

## schedule(dynamic, [n])

- Iteration of loop are divided into chunks containing n iterations each.



```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <omp.h>
4 int main(int argc, char* argv[])
5 {
6     omp_set_num_threads(2);
7     printf("With no chunksize passed:\n");
8     #pragma omp parallel for schedule(dynamic)
9     for(int i = 0; i < 10; i++)
10    {
11        printf("Thread %d processes iteration %d.\n", omp_get_thread_num(), i);
12    }
13    printf("With a chunksize of 2:\n");
14    #pragma omp parallel for schedule(dynamic, 2)
15    for(int i = 0; i < 10; i++)
16    {
17        printf("Thread %d processes iteration %d.\n", omp_get_thread_num(), i);
18    }
19    return 0;
20 }
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>
$./<file_name>
```

- Default chunk size is 1.

\$ ./parallel\_for\_loop

With no chunksize passed:

Thread 0 processes iteration 0.  
Thread 0 processes iteration 2.  
Thread 0 processes iteration 3.  
Thread 0 processes iteration 4.  
Thread 0 processes iteration 5.  
Thread 0 processes iteration 6.  
Thread 0 processes iteration 7.  
Thread 0 processes iteration 8.  
Thread 0 processes iteration 9.  
Thread 1 processes iteration 1.

With a chunksize of 2:

Thread 1 processes iteration 0.  
Thread 1 processes iteration 1.  
Thread 1 processes iteration 4.  
Thread 1 processes iteration 5.  
Thread 1 processes iteration 6.  
Thread 1 processes iteration 7.  
Thread 1 processes iteration 8.  
Thread 1 processes iteration 9.  
Thread 0 processes iteration 2.  
Thread 0 processes iteration 3.

## SCHEDULE Clause.. cont

### **schedule(guided, [n])**

- If you specify n, that is the minimum chunk size that each thread should posses..
- Size of each successive chunk is exponentially decreasing.
- Initial chunk size  
 $\max(\text{number\_of\_iterations} / \text{number\_of\_threads}, n)$   
Subsequent chunk consist of  
 $\max(\text{remaining\_iterations}/\text{number\_of\_threads}, n)$  iterations

Schedule(runtime):export OMP\_SCHEDULE “STATIC,4”

- Determine the scheduling type at run time by the OMP\_SCHEDULE environment variable.

## For loop - SCHEDULE(guided)



```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <omp.h>
4
5 int main(int argc, char* argv[])
6 {
7     omp_set_num_threads(4);
8     #pragma omp parallel for schedule(guided)
9     for(int i = 0; i < 10; i++)
10    {
11        printf("Thread %d processes iteration %d.\n", omp_get_thread_num(), i);
12    }
13
14    return 0;
15 }
```

```
$ ./parallel_loop_scheduling
Thread 2 processes iteration 7.
Thread 2 processes iteration 8.
Thread 2 processes iteration 9.
Thread 1 processes iteration 5.
Thread 1 processes iteration 6.
Thread 3 processes iteration 3.
Thread 3 processes iteration 4.
Thread 0 processes iteration 0.
Thread 0 processes iteration 1.
Thread 0 processes iteration 2.
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>
$./<file_name>
```

## For loop - SCHEDULE(runtime)

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <omp.h>
4
5 int main(int argc, char* argv[])
6 {
7     omp_set_num_threads(4);
8     #pragma omp parallel for schedule(runtime)
9     for(int i = 0; i < 10; i++)
10    {
11        printf("Thread %d processes iteration %d.\n", omp_get_thread_num(), i);
12    }
13
14    return 0;
15 }
```

```
./parallel_for_runtime
Thread 0 processes iteration 0.
Thread 0 processes iteration 4.
Thread 0 processes iteration 5.
Thread 0 processes iteration 6.
Thread 0 processes iteration 7.
Thread 0 processes iteration 8.
Thread 1 processes iteration 3.
Thread 2 processes iteration 2.
Thread 0 processes iteration 9.
Thread 3 processes iteration 1.
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>
$./<file_name>
```

## SECTIONS Construct

- The sections construct is a non-iterative worksharing construct that contains a set of structured blocks.
- The blocks are distributed among and executed by the threads in a team.
- Each structured block is executed once by one of the threads in the team in the context of its implicit task.

The parallel directive supports the following clauses:

- **private(list)**
- **firstprivate(list)**
- **lastprivate([modifier :] list)**
- **reduction([modifier,] identifier : list)**
- **nowait**

# SECTIONS - PRIVATE

```
● ● ●  
1 #include<stdio.h>  
2 #include <omp.h>  
3 #define N      5  
4 int main () {  
5     int i;  
6     float a[N], b[N], c[N], d[N],e[N];  
7     /* Some initializations */  
8     for (i=0; i < N; i++) {  
9         a[i] = i * 1.5;  
10        b[i] = i + 22.35;  
11    }  
12    /* Fork a team of threads with each thread having a private i variable and shared varable a,b,c,chuk */  
13    #pragma omp parallel shared(a,b,c,d,e) private(i)  
14    {  
15        #pragma omp sections  
16        {  
17            #pragma omp section  
18            for (i=0; i < N; i++){  
19                c[i] = a[i] + b[i];  
20                printf("section 1 # Working thread : %d | %f + %f = %f\n",omp_get_thread_num(),a[i],b[i],c[i]);  
21            }  
22            #pragma omp section  
23            for (i=0; i < N; i++){  
24                d[i] = a[i] * b[i];  
25                printf("section 2 # Working thread : %d | %f * %f = %f\n",omp_get_thread_num(),a[i],b[i],d[i]);  
26            }  
27        } /* end of sections */  
28    } /* end of parallel section */  
29    return 0;  
30 }  
31
```

./sections

```
section 2 # Working thread : 2 | 0.000000 * 22.350000 = 0.000000  
section 2 # Working thread : 2 | 1.500000 * 23.350000 = 35.025002  
section 2 # Working thread : 2 | 3.000000 * 24.350000 = 73.050003  
section 2 # Working thread : 2 | 4.500000 * 25.350000 = 114.075005  
section 2 # Working thread : 2 | 6.000000 * 26.350000 = 158.100006  
section 1 # Working thread : 0 | 0.000000 + 22.350000 = 22.350000  
section 1 # Working thread : 0 | 1.500000 + 23.350000 = 24.850000  
section 1 # Working thread : 0 | 3.000000 + 24.350000 = 27.350000  
section 1 # Working thread : 0 | 4.500000 + 25.350000 = 29.850000  
section 1 # Working thread : 0 | 6.000000 + 26.350000 = 32.349998
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

# SECTIONS - FIRSTPRIVATE



```
1 #include<stdio.h>
2 #include <omp.h>
3 #define N      5
4 int main () {
5     int i;
6     float a[N], b[N], c[N], d[N], e[N];
7     /* Some initializations */
8     for (i=0; i < N; i++) {
9         a[i] = i * 1.5;
10        b[i] = i + 22.35;
11    }
12    /* Fork a team of threads with each thread having a private i variable and shared varable a,b,c,chuk */
13    #pragma omp parallel shared(a,b,c,d,e) firstprivate(i)
14    {
15        #pragma omp sections nowait
16        {
17            #pragma omp section
18            for (i=0; i < N; i++){
19                c[i] = a[i] + b[i];
20                printf("section 1 # Working thread : %d | %f + %f = %f\n",omp_get_thread_num(),a[i],b[i],c[i]);
21            }
22            #pragma omp section
23            for (i=0; i < N; i++){
24                d[i] = a[i] * b[i];
25                printf("section 2 # Working thread : %d | %f * %f = %f\n",omp_get_thread_num(),a[i],b[i],d[i]);
26            }
27        } /* end of sections */
28    } /* end of parallel section */
29    return 0;
30 }
```

\$ ./sections

```
section 2 # Working thread : 2 | 0.000000 * 22.350000 = 0.000000
section 2 # Working thread : 2 | 1.500000 * 23.350000 = 35.025002
section 2 # Working thread : 2 | 3.000000 * 24.350000 = 73.050003
section 2 # Working thread : 2 | 4.500000 * 25.350000 = 114.075005
section 2 # Working thread : 2 | 6.000000 * 26.350000 = 158.100006
section 1 # Working thread : 0 | 0.000000 + 22.350000 = 22.350000
section 1 # Working thread : 0 | 1.500000 + 23.350000 = 24.850000
section 1 # Working thread : 0 | 3.000000 + 24.350000 = 27.350000
section 1 # Working thread : 0 | 4.500000 + 25.350000 = 29.850000
section 1 # Working thread : 0 | 6.000000 + 26.350000 = 32.349998
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>
$./<file_name>
```

# SECTIONS - LASTPRIVATE

```
1 #include<stdio.h>
2 #include <omp.h>
3 #define N      5
4 int main () {
5     int i;
6     float a[N], b[N], c[N], d[N], e[N];
7     /* Some initializations */
8     for (i=0; i < N; i++) {
9         a[i] = i * 1.5;
10        b[i] = i + 22.35;
11    }
12    /* Fork a team of threads with each thread having
13    #pragma omp parallel shared(a,b,c,d,e)
14    {
15        #pragma omp sections nowait lastprivate(i)
16        {
17            #pragma omp section
18            for (i=0; i < N; i++){
19                c[i] = a[i] + b[i];
20                printf("section 1 # Working thread : %d | %f + %f = %f\n",omp_get_thread_num(),a[i],b[i],c[i]);
21            }
22            #pragma omp section
23            for (i=0; i < N; i++){
24                d[i] = a[i] * b[i];
25                printf("section 2 # Working thread : %d | %f * %f = %f\n",omp_get_thread_num(),a[i],b[i],d[i]);
26            }
27        } /* end of sections */
28    } /* end of parallel section */
29    return 0;
}
```

```
$ ./sections
section 1 # Working thread : 0 | 0.000000 + 22.350000 =
22.350000
section 1 # Working thread : 0 | 1.500000 + 23.350000 =
24.850000
section 1 # Working thread : 0 | 3.000000 + 24.350000 =
27.350000
section 1 # Working thread : 0 | 4.500000 + 25.350000 =
29.850000
section 1 # Working thread : 0 | 6.000000 + 26.350000 =
32.349998
section 2 # Working thread : 1 | 0.000000 * 22.350000 =
0.000000
section 2 # Working thread : 1 | 1.500000 * 23.350000 =
35.025002
section 2 # Working thread : 1 | 3.000000 * 24.350000 =
73.050003
How to compile and run?
$gcc -fopenmp <file_name.c>-O<file_name>
$./<file_name>
section 2 # Working thread : 1 | 4.500000 * 25.350000 =
114.075005
section 2 # Working thread : 1 | 6.000000 * 26.350000 =
158.100006
```

# SECTIONS - REDUCTION

```
● ● ●  
1 #include <stdio.h>  
2 #include <omp.h>  
3 #define NUM 100  
4 int isprime( int x )  
5 {  
6     for( int y = 2; y * y <= x; y++ )  
7     {  
8         if( x % y == 0 )  
9             return 0;  
10    }  
11    return 1;  
12 }  
13 int main( )  
14 {  
15     int sum = 0;  
16     int i;  
17 #pragma omp parallel  
18 {  
19     #pragma omp sections reduction (+:sum)  
20     {  
21         #pragma omp section  
22         {  
23             for( int i = 2; i <= NUM ; i++ )  
24             {  
25                 sum += isprime ( i );  
26                 }  
27             printf( "Number of primes numbers: %d\n", sum );  
28         }  
29         #pragma omp section  
30         {  
31             for( int i = 2; i <= NUM ; i++ )  
32             {  
33                 sum += isprime ( i );  
34                 }  
35             printf( "Number of primes numbers: %d\n", sum );  
36         }  
37     }  
38 }  
39 }  
40     return 0;  
41 }
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

\$ ./reduction

Number of primes numbers: 25

Number of primes numbers: 25

# SECTIONS - NOWAIT

```
● ● ●  
1 #include<stdio.h>  
2 #include <omp.h>  
3 #define N      5  
4 int main () {  
5     int i;  
6     float a[N], b[N], c[N], d[N],e[N];  
7     /* Some initializations */  
8     for (i=0; i < N; i++) {  
9         a[i] = i * 1.5;  
10        b[i] = i + 22.35;  
11    }  
12    /* Fork a team of threads with each thread having a private i variable and shared varable a,b,c,chuk */  
13    #pragma omp parallel shared(a,b,c,d,e) private(i)  
14    {  
15        #pragma omp sections nowait  
16        {  
17            #pragma omp section  
18            for (i=0; i < N; i++){  
19                c[i] = a[i] + b[i];  
20                printf("section 1 # Working thread : %d | %f + %f = %f\n",omp_get_thread_num(),a[i],b[i],c[i]);  
21            }  
22            #pragma omp section  
23            for (i=0; i < N; i++){  
24                d[i] = a[i] * b[i];  
25                printf("section 2 # Working thread : %d | %f * %f = %f\n",omp_get_thread_num(),a[i],b[i],d[i]);  
26            }  
27        } /* end of sections */  
28    } /* end of parallel section */  
29    return 0;  
30 }  
31
```

\$ ./sections

```
section 2 # Working thread : 1 | 0.000000 * 22.350000 = 0.000000  
section 2 # Working thread : 1 | 1.500000 * 23.350000 = 35.025002  
section 2 # Working thread : 1 | 3.000000 * 24.350000 = 73.050003  
section 2 # Working thread : 1 | 4.500000 * 25.350000 = 114.075005  
section 2 # Working thread : 1 | 6.000000 * 26.350000 = 158.100006  
section 1 # Working thread : 0 | 0.000000 + 22.350000 = 22.350000  
section 1 # Working thread : 0 | 1.500000 + 23.350000 = 24.850000  
section 1 # Working thread : 0 | 3.000000 + 24.350000 = 27.350000  
section 1 # Working thread : 0 | 4.500000 + 25.350000 = 29.850000  
section 1 # Working thread : 0 | 6.000000 + 26.350000 = 32.349998
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>  
$./<file_name>
```

## SINGLE Construct

- single is a clause that must be used in a parallel region; it tells OpenMP that the associated block must be executed by one thread only, .

**The clause to appear on the single construct, which is one of the following:**

- **private(list)**
- **firstprivate(list)**
- **copyprivate(list)**
- **nowait**
- master will be executed by the master only while single can be executed by whichever thread reaching the region first.

# SINGLE-PRIVATE



```
1 #include<stdio.h>
2 #include <omp.h>
3 #define N      5
4 int main () {
5     int i;
6     float a[N], b[N], c[N], d[N];
7     /* Some initializations */
8     for (i=0; i < N; i++) {
9         a[i] = i * 1.5;
10        b[i] = i + 22.35;
11    }
12    /* Fork a team of threads with each thread having a private i variable and shared varable a,b,c,d */
13 #pragma omp parallel shared(a,b,c,d) private(i)
14 {
15     #pragma omp single
16
17     {
18         for (i=0; i < N; i++){
19             c[i] = a[i] + b[i];
20             printf("1st loop # Working thread : %d # %f + %f = %f\n",omp_get_thread_num(),a[i],b[i],c[i]);
21         }
22         for (i=0; i < N; i++){
23             d[i] = a[i] * b[i];
24             printf("2nd loop # Working thread : %d # %f * %f = %f\n",omp_get_thread_num(),a[i],b[i],d[i]);
25         }
26     }
27     /* end of sections */
28 } /* end of parallel section */
29 return 0;
30 }
```

\$ ./single

```
1st loop # Working thread : 0 # 0.000000 + 22.350000 = 22.350000
1st loop # Working thread : 0 # 1.500000 + 23.350000 = 24.850000
1st loop # Working thread : 0 # 3.000000 + 24.350000 = 27.350000
1st loop # Working thread : 0 # 4.500000 + 25.350000 = 29.850000
1st loop # Working thread : 0 # 6.000000 + 26.350000 = 32.349998
2nd loop # Working thread : 0 # 0.000000 * 22.350000 = 0.000000
2nd loop # Working thread : 0 # 1.500000 * 23.350000 = 35.025002
2nd loop # Working thread : 0 # 3.000000 * 24.350000 = 73.050003
2nd loop # Working thread : 0 # 4.500000 * 25.350000 = 114.075005
2nd loop # Working thread : 0 # 6.000000 * 26.350000 = 158.100006
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>
$./<file_name>
```

# SINGLE-FIRSTPRIVATE

```
● ● ●  
1 #include<stdio.h>  
2 #include <omp.h>  
3 #define N 5  
4 int main () {  
5     int i = 1;  
6     float a[N], b[N], c[N], d[N];  
7     /* Some initializations */  
8     for (i=0; i < N; i++) {  
9         a[i] = i * 1.5;  
10        b[i] = i + 22.35;  
11    }  
12    /* Fork a team of threads with each thread having a  
13 #pragma omp parallel shared(a,b,c,d) firstprivate(i)  
14    {  
15        #pragma omp single  
16        {  
17            for (i=0; i < N; i++){  
18                c[i] = a[i] + b[i];  
19                printf("1st loop # Working thread : %d # %f + %f = %f\n",omp_get_thread_num(),a[i],b[i],c[i]);  
20            }  
21            for (i=0; i < N; i++){  
22                d[i] = a[i] * b[i];  
23                printf("2nd loop # Working thread : %d # %f * %f = %f\n",omp_get_thread_num(),a[i],b[i],d[i]);  
24            }  
25        }  
26    }  
27    /* end of sections */  
28 } /* end of parallel section */  
29 return 0;  
30 }
```

```
$ ./single  
1st loop # Working thread : 0 # 0.000000 + 22.350000 =  
22.350000  
1st loop # Working thread : 0 # 1.500000 + 23.350000 =  
24.850000  
1st loop # Working thread : 0 # 3.000000 + 24.350000 =  
27.350000  
1st loop # Working thread : 0 # 4.500000 + 25.350000 =  
29.850000  
1st loop # Working thread : 0 # 6.000000 + 26.350000 =  
32.349998  
2nd loop # Working thread : 0 # 0.000000 * 22.350000 =  
0.000000  
2nd loop # Working thread : 0 # 1.500000 * 23.350000 =  
35.025002  
2nd loop # Working thread : 0 # 3.000000 * 24.350000 =  
73.050003  
How to compile and run?  
$gcc -fopenmp <file_name.c> -o <file_name>  
2nd loop # Working thread : 0 # 4.500000 * 25.350000 =  
114.075005  
2nd loop # Working thread : 0 # 6.000000 * 26.350000 =  
158.100006
```

# SINGLE-COPYPRIVATE

```
● ● ●  
1 #include <stdio.h>  
2 #include <stdlib.h>  
3 #include <omp.h>  
4 int main(int argc, char* argv[]){  
5 {  
6     int a = 123;  
7  
8     #pragma omp parallel default(none) firstprivate(a)  
9     {  
10        printf("Thread %d: a = %d.\n", omp_get_thread_num(), a);  
11  
12        #pragma omp barrier  
13  
14        #pragma omp single copyprivate(a)  
15        {  
16            a = 456;  
17            printf("Thread %d executes the single construct and changes a to %d.\n", omp_get_thread_num(), a);  
18        }  
19  
20        printf("Thread %d: a = %d.\n", omp_get_thread_num(), a);  
21    }  
22  
23    return 0;  
24 }
```

\$ ./mp\_copyprivate

Thread 0: a = 123.

Thread 1: a = 123.

Thread 2: a = 123.

Thread 3: a = 123.

Thread 0 executes the single construct and changes a to 456.

Thread 2: a = 456.

Thread 3: a = 456.

Thread 1: a = 456.

Thread 0: a = 456.

How to compile and run?

\$gcc -fopenmp <file\_name.c> -o <file\_name>

\$.</file\_name>

# SINGLE-NOWAIT

```
1 #include<stdio.h>
2 #include <omp.h>
3 #define N      5
4 int main () {
5     int i;
6     float a[N], b[N], c[N], d[N];
7     /* Some initializations */
8     for (i=0; i < N; i++) {
9         a[i] = i * 1.5;
10        b[i] = i + 22.35;
11    }
12    /* Fork a team of threads with each thread having a private i variable and shared varable a,b,c,d */
13 #pragma omp parallel shared(a,b,c,d) private(i)
14 {
15     #pragma omp single nowait
16
17     {
18         for (i=0; i < N; i++){
19             c[i] = a[i] + b[i];
20             printf("1st loop # Working thread : %d # %f + %f = %f\n",omp_get_thread_num(),a[i],b[i],c[i]);
21         }
22         for (i=0; i < N; i++){
23             d[i] = a[i] * b[i];
24             printf("2nd loop # Working thread : %d # %f * %f = %f\n",omp_get_thread_num(),a[i],b[i],d[i]);
25         }
26     }
27     /* end of sections */
28 } /* end of parallel section */
29 return 0;
30 }
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o <file_name>
$./<file_name>
```

```
$ ./single
```

```
1st loop # Working thread : 0 # 0.000000 + 22.350000
= 22.350000
1st loop # Working thread : 0 # 1.500000 + 23.350000
= 24.850000
1st loop # Working thread : 0 # 3.000000 + 24.350000
= 27.350000
1st loop # Working thread : 0 # 4.500000 + 25.350000
= 29.850000
1st loop # Working thread : 0 # 6.000000 + 26.350000
= 32.349998
2nd loop # Working thread : 0 # 0.000000 *
22.350000 = 0.000000
2nd loop # Working thread : 0 # 1.500000 *
23.350000 = 35.025002
2nd loop # Working thread : 0 # 3.000000 *
24.350000 = 73.050003
2nd loop # Working thread : 0 # 4.500000 *
25.350000 = 114.075005
2nd loop # Working thread : 0 # 6.000000 *
26.350000 = 158.100006
```

# Questions for Thought

- What happens if the number of threads and the number of SECTIONs are different? What if there are more threads than SECTIONs? fewer threads than SECTIONs?
- Which thread executes which SECTION?

# Synchronization Constructs

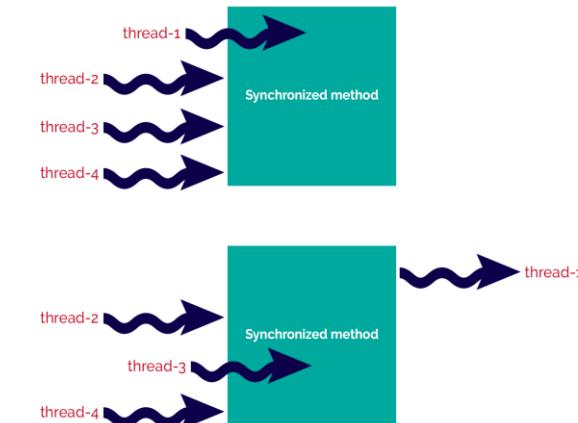
- Motivation: Consider a simple example where two threads on two different processors are both trying to increment a variable  $x$  at the same time (assume  $x$  is initially 0).

```
THREAD 1:  
increment(x)  
{  
    x = x + 1;  
}  
THREAD 1:
```

```
10 LOAD A, (x address)  
20 ADD A, 1  
30 STORE A, (x address)
```

```
THREAD 2:  
increment(x)  
{  
    x = x + 1;  
}  
THREAD 2:
```

```
10 LOAD A, (x address)  
20 ADD A, 1  
30 STORE A, (x address)
```



- The incrementation of  $x$  must be synchronized between the two threads to insure that the correct result is produced.
- OpenMP provides a variety of synchronization constructs that control how the execution of each thread proceeds relative to other team threads.

# Critical Construct

The critical construct, which is used inside parallel regions, tells OpenMP that the associated block is to be executed by every thread but no more than one thread at a time.

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

```
./critical
```

```
[Thread 0] Total before I add my value (1): 0.  
[Thread 0] Total after I added my value: 1.  
[Thread 3] Total before I add my value (6): 1.  
[Thread 3] Total after I added my value: 7.  
[Thread 2] Total before I add my value (2): 7.  
[Thread 2] Total after I added my value: 9.  
[Thread 1] Total before I add my value (1): 9.  
[Thread 1] Total after I added my value: 10.
```

# Critical Construct

```
int main(int argc, char* argv[])
{
    // Use 4 threads when creating OpenMP parallel regions
    omp_set_num_threads(4);

    int total = 0;

    // Create the parallel region
    #pragma omp parallel default(none) shared(total)
    {
        // Calculate my factorial
        int my_value = 1;
        for(int i = 2; i <= omp_get_thread_num(); i++)
        {
            my_value *= i;
        }

        // Add my value to the total
        #pragma omp critical
        {
            printf("[Thread %d] Total before I add my value (%d): %d.\n", omp_get_thread_num(), my_value, total);
            total += my_value;
            printf("[Thread %d] Total after I added my value: %d.\n", omp_get_thread_num(), total);
        }
    }

    return 0;
}
```

# Master Construct

master is a clause that must be used in a parallel region; it tells OpenMP that the associated block must be executed by the master thread only. The other threads do not wait at the end of the associated block as if there was an implicit barrier.

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

\$ ./Master

[Thread 0] Every thread executes this printf.  
[Thread 3] Every thread executes this printf.  
[Thread 1] Every thread executes this printf.  
[Thread 2] Every thread executes this printf.  
[Thread 0] Only the master thread executes this printf, which is me.

# Master Construct

```
int main(int argc, char* argv[])
{
    // Use 4 threads when creating OpenMP parallel regions
    omp_set_num_threads(4);

    // Create the parallel region
    #pragma omp parallel
    {
        printf("[Thread %d] Every thread executes this printf.\n", omp_get_thread_num());

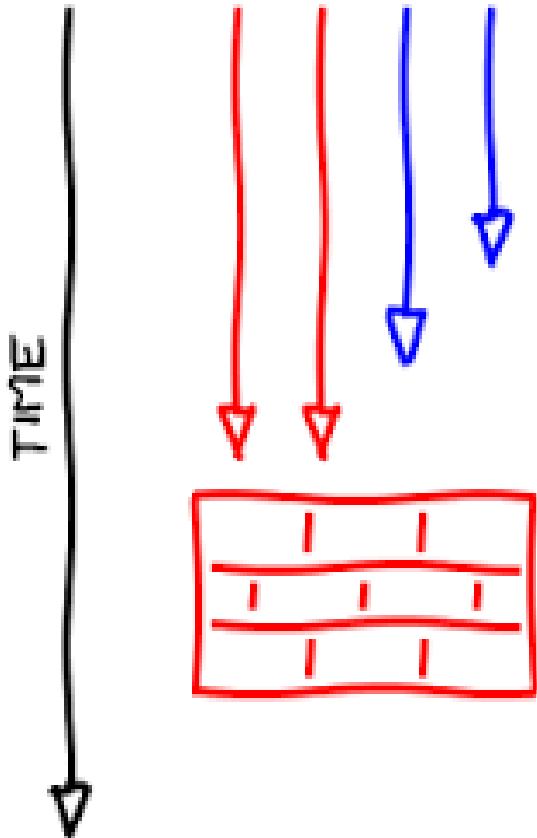
        #pragma omp barrier

        #pragma omp master
        {
            printf("[Thread %d] Only the master thread executes this printf, which is me.\n", omp_get_thread_num());
        }
    }

    return 0;
}
```

# BARRIER Construct

- Explicit wait for other threads to Complete their task



How to compile and run?

```
$gcc -fopenmp <file_name.c> -o  
<file_name>  
$./<file_name>
```

```
$ ./barrier
```

```
[Thread 0] I print my first message.  
[Thread 3] I print my first message.  
[Thread 1] I print my first message.  
[Thread 2] I print my first message.  
The barrier is complete, which means all  
threads have printed their first message.  
[Thread 3] I print my second message.  
[Thread 2] I print my second message.  
[Thread 1] I print my second message.  
[Thread 0] I print my second message.
```

# BARRIER



```
1 int main(int argc, char* argv[])
2 {
3     // Use 4 threads when we create a parallel region
4     omp_set_num_threads(4);
5
6     // Create the parallel region
7     #pragma omp parallel
8     {
9         // Threads print their first message
10        printf("[Thread %d] I print my first message.\n", omp_get_thread_num());
11
12        // Make sure all threads have printed their first message before moving on.
13        #pragma omp barrier
14
15        // One thread indicates that the barrier is complete.
16        #pragma omp single
17        {
18            printf("The barrier is complete, which means all threads have printed their first message.\n");
19        }
20
21        // Threads print their second message
22        printf("[Thread %d] I print my second message.\n", omp_get_thread_num());
23    }
24
25    return 0;
26 }
27
```

# ATOMIC

- The atomic construct ensures that a specific storage location is accessed atomically, rather than exposing it to the possibility of multiple, simultaneous reading and writing threads that may result in indeterminate values.

```
● ● ●  
1 int main(int argc, char* argv[])
2 {
3     // Use 4 threads when creating OpenMP parallel
4     omp_set_num_threads(4);
5     int total = 0;
6     // Create the parallel region
7     #pragma omp parallel default(none) shared(total)
8     {
9         for(int i = 0; i < 10; i++)
10        {
11            // Atomically add one to the total
12            #pragma omp atomic
13            total++;
14        }
15    }
16    printf("Total = %d.\n", total);
17    return 0;
18 }
```

How to compile and run?

```
$gcc -fopenmp <file_name.c> -o
<file_name>
$./<file_name>
```

```
$ ./atomic
Total = 40.
```

# FLUSH Directive

- Identifies a synchronization point at which the implementation must provide a consistent view of memory
- Thread-visible variables are written back to memory at this point.
- Necessary to instruct the compiler that a variable must be written to/read from the memory system, i.e. that a variable cannot be kept in a local CPU register
  - Keeping a variable in a register in a loop is very common when producing efficient machine language code for a loop

C/C++: #pragma omp flush (list) newline

# FLUSH Directive (cont.)

- The optional list contains a list of named variables that will be flushed in order to avoid flushing all variables. For pointers in the list, the pointer itself is flushed, not the object to which it points.
- Implementations must ensure any prior modifications to thread-visible variables are visible to all threads after this point; i.e., compilers must restore values from registers to memory, hardware might need to flush write buffers, etc.
- The FLUSH directive is implied for the directives shown in the table below. The directive is not implied if a NOWAIT clause is present

# FLUSH Directive (cont.)

- The FLUSH directive is implied for the directives shown in the table below. The directive is not implied if a NOWAIT clause is present.

| C/C++                          |
|--------------------------------|
| barrier                        |
| parallel – upon entry and exit |
| critical – upon entry and exit |
| ordered – upon entry and exit  |
| for – upon exit                |
| sections – upon exit           |
| single – upon exit             |

# ORDERED Directive

- Specifies that iterations of the enclosed loop will be executed in the same order as if they were executed on a serial processor
- Threads will need to wait before executing their chunk of iterations if previous iterations haven't completed yet.
- Used within a DO / for loop with an ORDERED clause
- The ORDERED directive provides a way to "fine tune" where ordering is to be applied within a loop. Otherwise, it is not required.
- An ORDERED directive can only appear in the dynamic extent of the following directives:
  - DO or PARALLEL DO (Fortran)
  - for or parallel for (C/C++)
- Only one thread is allowed in an ordered section at any time
- It is illegal to branch into or out of an ORDERED block.
- An iteration of a loop must not execute the same ORDERED directive more than once, and it must not execute more than one ORDERED directive.
- A loop that contains an ORDERED directive must be a loop with an ORDERED clause.

# ORDERED Example

# Data Scope Attribute Clauses

- Also called data sharing attribute clauses
- Because OpenMP is based upon the shared memory programming model, most variables are shared by default.
- Global variables include:
  - C: File scope variables, static
- Private variables include:
  - Loop index variables
- Clauses used to explicitly define how variables should be scoped
- include:
  - PRIVATE
  - FIRSTPRIVATE
  - LASTPRIVATE
  - SHARED
  - DEFAULT
  - REDUCTION
  - COPYIN

# Data Scope Attribute Clauses (cont.)

- Used in conjunction with several directives (PARALLEL, DO/for, and SECTIONS) to control the scoping of enclosed variables.
- Provide the ability to control the data environment during execution of parallel constructs
  - Define how and which data variables in the serial section of the program are transferred to the parallel sections of the program (and back)
  - Define which variables will be visible to all threads in the parallel sections and which variables will be privately allocated to all threads
- Effective only within their lexical/static extent

# PRIVATE and SHARED Clauses

- PRIVATE Clause
  - Declares variables in its list to be private to each thread
  - A new object of the same type is declared once for each thread in the team.
  - All references to the original object are replaced with references to the new object.
  - Variables declared PRIVATE should be assumed to be uninitialized for each thread.
- SHARED Clause
  - Declares variables in its list to be shared among all threads in the team
  - A shared variable exists in only one memory location and all threads can read or write to that address.
  - It is the programmer's responsibility to ensure that multiple threads properly access SHARED variables (such as via CRITICAL sections)

# FIRSTPRIVATE and LASTPRIVATE Clauses

- FIRSTPRIVATE Clause
  - Combines the behaviour of the PRIVATE clause with automatic initialization of the variables in its list
  - Listed variables are initialized according to the value of their original objects prior to entry into the parallel or work-sharing construct.
- LASTPRIVATE Clause
  - Combines the behaviour of the PRIVATE clause with a copy from the last loop iteration or section to the original variable object
  - The value copied back into the original variable object is obtained from the last (sequentially) iteration or section of the enclosing construct.
  - For example, the team member that executes the final iteration for a DO section, or the team member that executes the last SECTION of a SECTIONS context, performs the copy with its own values.

# DEFAULT Clause

- Allows the user to specify a default scope for all variables in the lexical extent of any parallel region
- Specific variables can be exempted from the default using the PRIVATE, SHARED, FIRSTPRIVATE, LASTPRIVATE, and REDUCTION clauses.
- The C/C++ OpenMP specification does not include private or firstprivate as a possible default. However, actual implementations may provide this option.
- Using NONE as a default requires that the programmer explicitly scope all variables.

# REDUCTION Clause

- Performs a reduction on the variables that appear in its list.
- A private copy for each list variable is created for each thread. At the end of the reduction, the reduction operator is applied to all private copies of the shared variable, and the final result is written to the global shared variable.
- Syntax:  
reduction (operator: list)

# REDUCTION Clause Example

# THREADPRIVATE Directive

- Used to make global file scope variables (C/C++) or common blocks (Fortran) local and persistent to a thread through the execution of multiple parallel regions
- Must appear after the declaration of listed variables/common blocks
- Each thread then gets its own copy of the variable/common block, so data written by one thread is not visible to other threads.
- On first entry to a parallel region, data in THREADPRIVATE variables and common blocks should be assumed undefined, unless a COPYIN clause is specified in the PARALLEL directive
- THREADPRIVATE variables differ from PRIVATE variables because they are able to persist between different parallel sections of a code.
- Data in THREADPRIVATE objects is guaranteed to persist only if the dynamic threads mechanism is "turned off" and the number of threads in different parallel regions remains constant. The default setting of dynamic threads is undefined.

# THREADPRIVATE Example

# COPYIN Clause

- Provides a means for assigning the same value to THREADPRIVATE variables for all threads in the team
- List contains the names of variables to copy. In Fortran, the list can contain both the names of common blocks and named variables.
- The master thread variable is used as the copy source. The team threads are initialized with its value upon entry into the parallel construct

# Run-time Library Routines

- The OpenMP standard defines an API for library calls that perform a variety of functions:
  - Query the number of threads/processors, set number of threads to use
  - General purpose locking routines (semaphores)
  - Portable wall clock timing routines
  - Set execution environment functions: nested parallelism, dynamic adjustment of threads.
  - For C/C++, it may be necessary to specify the include file "omp.h".

Note: Your implementation may or may not support nested parallelism and/or dynamic threads. If nested parallelism is supported, it is often only nominal, in that a nested parallel region may only have one thread.

# OpenMP Environment Variables

- To set number of threads during execution

```
export OMP_NUM_THREADS=4
```

- To allow run time system to determine the number of threads

```
export OMP_DYNAMIC=TRUE
```

- To allow nesting of parallel region

```
export OMP_NESTED=TRUE
```